



Hanzo 3.0: AGI for All — Next-Generation Agent Management with Agent Chaining and Vector Database Architecture

Zach Kelling

Hanzo AI

2023

Hanzo 3.0: AGI for All — Next-Generation Agent Management with Agent Chaining and Vector Database Architecture

Zach Kelling
CTO, Hanzo AI (Techstars '17)
zach@hanzo.ai

2023

Abstract

We present Hanzo 3.0, a platform for managing multiple AI agents at scale. The system introduces *Agent Chaining*, an architecture for composing open-source AI protocols into pipelines; a *Vector Database Backend* for infinite-recall retrieval-augmented generation; *4-bit quantization* for efficient local and private model deployment; and *multidimensional user interfaces* spanning voice, text, desktop, mobile, wearable, and AR/VR modalities. We describe the platform’s architecture, its applications across finance, healthcare, manufacturing, and other industries, and its implications for the trajectory toward artificial general intelligence.

1 Introduction

As AI capabilities expand, organizations increasingly deploy not one model but constellations of specialized agents—each handling a distinct task, domain, or modality. Managing these agent fleets introduces operational challenges: monitoring performance across agents, distributing tasks equitably, maintaining consistent configuration, and providing a unified interface to end users who neither know nor care which agent is responding.

Hanzo 3.0 addresses these challenges with a platform purpose-built for multi-agent management. The system treats AI agents as first-class operational entities with lifecycle management, versioning, monitoring, and collaboration primitives.

2 Architecture

2.1 Agent Chaining

Agent Chaining is the core abstraction. Rather than deploying agents as isolated services, Hanzo composes them into directed acyclic graphs

(DAGs) where the output of one agent feeds the input of the next. This enables:

- **Protocol composition.** Open-source AI protocols (language models, vision encoders, speech recognizers) snap together via typed interfaces.
- **Conditional routing.** Chains branch based on confidence scores, input type, or domain classification.
- **Fallback cascades.** If a primary agent fails or returns low-confidence output, the chain automatically routes to a secondary.

2.2 Vector Database Backend

The platform stores all agent interactions, user context, and domain knowledge in a vector database optimized for high-concurrency, low-latency similarity search. This provides:

- **Infinite recall.** Agents retrieve relevant context from the full history of interactions, not just a fixed context window.
- **Retrieval-augmented generation (RAG).** Every agent query is enriched

with semantically similar prior interactions and domain documents.

- **Real-time indexing.** New interactions are indexed immediately, ensuring the knowledge base is always current.

2.3 4-Bit Quantization

Hanzo supports 4-bit model quantization, reducing model size by approximately $8\times$ compared to FP32 with minimal accuracy degradation. This enables:

- Local deployment on consumer hardware (laptops, phones).
- Private inference without sending data to cloud endpoints.
- Hybrid architectures where lightweight local models handle routine queries and cloud models handle complex ones.

2.4 Token Length Research

The platform incorporates ongoing research into extending effective context length. Longer token windows enable more detailed, contextually rich agent responses—critical for domains like legal analysis, medical records review, and financial modeling where relevant context spans thousands of tokens.

3 Multi-Agent Dashboard

The dashboard provides a unified operational view:

- **Real-time metrics.** Conversation count, response latency, accuracy, and throughput per agent.
- **Filtering and grouping.** Agents sorted by model type, domain, performance tier, or deployment location.
- **Configuration management.** Parameters and training data managed at individual or group level.

- **Versioning.** Save and restore agents at any version, enabling safe rollback.

- **Parallel training and deployment.** Train and deploy multiple agents concurrently; clone existing agents for rapid scaling.

4 Task Distribution

The Task Distribution System ensures equitable and efficient workload allocation:

- **Round-robin.** Simple rotation across available agents.
- **Least-busy.** Route to the agent with the lowest current load.
- **Performance-based.** Route to the agent with the highest accuracy or lowest latency for the given task type.

Automated scaling adjusts the active agent count in response to workload fluctuations.

5 Multidimensional Interfaces

Hanzo 3.0 supports interaction across every modality:

Modality	Capabilities
Voice	Context-aware conversation, commands
Text	Chat, structured queries, forms
Desktop/Mobile	Native applications, notifications
Watch/Wearable	Glanceable summaries, quick actions
AR contact lens	Ambient overlays, real-time data
HoloLens/VR	Spatial data exploration, collaboration

Table 1: Supported interface modalities.

A novel contribution is the treatment of *time as an interface dimension*. Users navigate their data timeline—reviewing past interactions, projecting future scenarios—transforming time from passive flow to active exploration tool. In 3D/4D environments, temporal navigation enables trend visualization and predictive modeling.

5.1 Ad Hoc AI-Generated Interfaces

The platform can construct personalized interfaces in real time based on user context and task requirements. These range from simple 2D displays to spatial 3D environments, generated on demand rather than pre-designed.

6 Multi-Agent Collaboration

For tasks requiring multiple perspectives, agents collaborate through:

- **Sequential handoff.** Agents take turns, each adding domain expertise.
- **Parallel consensus.** Multiple agents process the same input and the system synthesizes responses.
- **Feedback integration.** User feedback flows back through reinforcement learning, continuously improving agent behavior.

7 Industry Applications

7.1 Finance

A unified platform integrating bank accounts, crypto wallets, and investment portfolios. Automated trading executes across asset classes based on user-defined goals. Risk management algorithms balance portfolios for optimal risk-adjusted returns. Real-time insights track performance, identify risks, and surface opportunities.

7.2 Healthcare

AR-enhanced interfaces for clinicians overlay patient data and predictive models during procedures. Agent chains coordinate across imaging analysis, records review, and treatment recommendation.

7.3 Manufacturing

Real-time AR visualization of production lines highlights bottlenecks, maintenance needs, and

safety concerns. Predictive analytics forecast equipment failure before it occurs.

7.4 Mining and Commodities

Agent chains span the full value chain: deposit forecasting, equipment monitoring, refining optimization, logistics routing, market analysis, and provenance tracking via blockchain.

8 Human Agency

The platform is designed to augment rather than replace human decision-making. Users set parameters, manage training data, and provide feedback—ensuring AI adapts to human preferences rather than the reverse. Voice recognition broadens accessibility, and every interface is designed so that technology serves the user.

9 Toward Singularity

Hanzo 3.0 represents a step toward artificial general intelligence through automated decision-making, continuous learning via reinforcement feedback, and increasingly autonomous agent coordination. The platform’s ethical commitments—user privacy, data security, regulatory compliance—ensure that this trajectory respects human rights and societal norms.

10 Roadmap

Phase I. Advanced machine learning algorithms, expanded third-party integrations.

Phase II. Broader language support, new domain specializations.

Phase III. Integral everyday companion—AI assistance spanning all aspects of daily life.

11 Conclusion

Hanzo 3.0 provides the operational infrastructure for the multi-agent era. Agent Chaining composes specialized models into capable

pipelines. Vector database storage delivers infinite recall. Quantization enables private, local deployment. Multidimensional interfaces meet users in whatever modality suits the task. The result is a platform that makes advanced AI management accessible across industries while maintaining the ethical foundations necessary for responsible progress toward artificial general intelligence.

References

- [1] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, L. Kaiser, and I. Polosukhin. Attention is all you need. In *NeurIPS*, 2017.
- [2] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin, C. Zhang, S. Agarwal, K. Slama, A. Ray, et al. Training language models to follow instructions with human feedback. In *NeurIPS*, 2022.
- [3] H. Chase. LangChain: Building applications with LLMs through composability. <https://github.com/langchain-ai/langchain>, 2022.
- [4] G. Yang, R. Nakano, and S. Richards. Auto-GPT: An autonomous GPT-4 experiment. <https://github.com/Significant-Gravitas/Auto-GPT>, 2023.
- [5] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *NeurIPS*, 2020.
- [6] T. Dettmers, M. Lewis, Y. Belkada, and L. Zettlemoyer. GPT3.int8(): 8-bit matrix multiplication for transformers at scale. In *NeurIPS*, 2022.
- [7] H. Touvron, T. Lavril, G. Izacard, X. Martinet, M.-A. Lachaux, T. Lacroix, B. Rozière, N. Goyal, E. Hambro, F. Azhar, et al. LLaMA: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*, 2023.
- [8] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao. ReAct: Synergizing reasoning and acting in language models. In *ICLR*, 2023.