



Hanzo API Gateway: Unified Commerce API Architecture

Zach Kelling

Hanzo AI

February 2019

Abstract

We present the Hanzo API Gateway, a unified commerce API architecture supporting both REST and GraphQL interfaces. The gateway provides rate limiting, caching, authentication, and request routing for e-commerce operations. We formalize the API design principles, implement a token bucket rate limiter with distributed state, and introduce a novel caching strategy that achieves 94% cache hit rates for product data. The system handles 180,000 requests per second at P99 latency of 23ms. Production deployments serve 1,200+ merchants with 99.99% availability over 18 months.

1 Introduction

Modern e-commerce platforms require APIs that serve diverse clients: web applications, mobile apps, third-party integrations, and internal services. These clients have varying requirements for data granularity, latency, and throughput. A unified API gateway must balance these needs while maintaining security, consistency, and performance.

Hanzo API Gateway addresses these challenges through:

1. **Dual Protocol Support:** REST for simplicity, GraphQL for flexibility
2. **Intelligent Rate Limiting:** Per-client, per-endpoint quotas with burst handling
3. **Multi-Layer Caching:** Edge, gateway, and origin caching with coherent invalidation
4. **Unified Authentication:** API keys, OAuth 2.0, and JWT with fine-grained permissions

1.1 Contributions

This paper contributes:

- A formal model for unified REST/GraphQL API design
- A distributed token bucket rate limiter with sub-millisecond overhead
- A cache coherence protocol for e-commerce data with strong consistency guarantees
- Production validation at 180K requests/second with 99.99% availability

2 API Design Principles

2.1 Resource Model

The Hanzo API exposes e-commerce entities as resources:

Definition 2.1 (Commerce Resource). *A resource R is a tuple $(type, id, attributes, relationships)$ where:*

- *$type \in \{Product, Order, Customer, Cart, Collection, Discount\}$*
- *id : globally unique identifier*
- *$attributes$: key-value properties*
- *$relationships$: references to related resources*

2.2 REST Interface

REST endpoints follow a consistent pattern:

Listing 1: REST Endpoint Structure

```
1 # Collection endpoints
2 GET    /v1/{resource}          # List resources
3 POST   /v1/{resource}          # Create resource
4
5 # Instance endpoints
6 GET    /v1/{resource}/{id}     # Read resource
7 PUT    /v1/{resource}/{id}     # Replace resource
8 PATCH  /v1/{resource}/{id}     # Update resource
9 DELETE /v1/{resource}/{id}     # Delete resource
10
11 # Nested resources
12 GET    /v1/{resource}/{id}/{nested}
```

2.2.1 Query Parameters

List endpoints support filtering, sorting, and pagination:

$$query = \{fields, filter, sort, page, limit\} \quad (1)$$

Listing 2: REST Query Example

```
1 GET /v1/products?fields=id,name,price
2   &filter[category]=electronics
3   &filter[price.gte]=100
4   &sort=-created_at
5   &page=2&limit=50
```

2.3 GraphQL Interface

GraphQL provides flexible querying:

Listing 3: GraphQL Schema Excerpt

```
1 type Product {
2   id: ID!
3   name: String!
4   description: String
5   price: Money!
6   variants: [Variant!]!
7   collections: [Collection!]!
8   inventory: Inventory!
9 }
10
11 type Query {
12   product(id: ID!): Product
13   products(
14     filter: ProductFilter
15     first: Int
16     after: String
17   ): ProductConnection!
18 }
19
```

```

20 type Mutation {
21   createProduct(input: ProductInput!): Product!
22   updateProduct(id: ID!, input: ProductInput!): Product!
23 }

```

2.3.1 Query Complexity Analysis

To prevent resource exhaustion, we analyze query complexity:

Definition 2.2 (Query Complexity). *For GraphQL query Q , complexity is:*

$$complexity(Q) = \sum_{f \in fields(Q)} cost(f) \cdot multiplier(f) \quad (2)$$

where $cost(f)$ is the base cost of field f and $multiplier(f)$ accounts for list cardinality.

Algorithm 1 GraphQL Complexity Analysis

```

1: function COMPUTECOMPLEXITY( $node, multiplier$ )
2:    $total \leftarrow 0$ 
3:   for  $field \in node.selections$  do
4:      $cost \leftarrow baseCost(field.name)$ 
5:      $mult \leftarrow multiplier$ 
6:     if  $isList(field.type)$  then
7:        $limit \leftarrow getLimit(field.arguments)$ 
8:        $mult \leftarrow mult \times \min(limit, maxLimit)$ 
9:     end if
10:     $total \leftarrow total + cost \times mult$ 
11:    if  $hasSelections(field)$  then
12:       $total \leftarrow total + COMPUTECOMPLEXITY(field, mult)$ 
13:    end if
14:  end for
15:  return  $total$ 
16: end function

```

Queries exceeding complexity threshold $\theta_{max} = 10,000$ are rejected.

3 Gateway Architecture

3.1 Request Flow




Figure 1: Gateway request flow

Request processing stages:

1. **Edge:** TLS termination, DDoS protection, geographic routing

2. **Authentication:** Validate credentials, extract identity
3. **Rate Limiting:** Check and decrement quota
4. **Cache Lookup:** Return cached response if valid
5. **Routing:** Forward to appropriate backend service
6. **Response Processing:** Transform, cache, return

3.2 Authentication

3.2.1 API Key Authentication

Listing 4: API Key Header

```
1 Authorization: Bearer sk_live_abc123...
```

API keys encode:

- Key type: $\{sk_live, sk_test, pk_live, pk_test\}$
- Merchant identifier
- Permission scopes
- Expiration (optional)

3.2.2 OAuth 2.0 for Third-Party Apps

Listing 5: OAuth Token Request

```
1 POST /oauth/token
2 Content-Type: application/x-www-form-urlencoded
3
4 grant_type=authorization_code
5 &code=AUTH_CODE
6 &client_id=CLIENT_ID
7 &client_secret=CLIENT_SECRET
8 &redirect_uri=https://app.example.com/callback
```

3.2.3 Permission Model

Definition 3.1 (API Permission). *A permission p is a tuple $(resource, action, scope)$ where:*

- $resource \in \mathcal{R}$: resource type
- $action \in \{read, write, delete\}$
- $scope \in \{own, all\}$: ownership constraint

Permission check:

$$authorized(key, request) = \exists p \in permissions(key) : matches(p, request) \quad (3)$$

4 Rate Limiting

4.1 Token Bucket Algorithm

We implement rate limiting using a distributed token bucket:

Definition 4.1 (Token Bucket). *A token bucket B is characterized by:*

- r : token replenishment rate (tokens/second)
- b : bucket capacity (max tokens)
- $tokens(t)$: current token count at time t

Token count evolution:

$$tokens(t) = \min(b, tokens(t_0) + r \cdot (t - t_0)) \quad (4)$$

Algorithm 2 Distributed Token Bucket

```
1: function CHECKRATELIMIT( $key, cost$ )
2:    $bucket \leftarrow$  GETBUCKET( $key$ )
3:    $now \leftarrow$  CURRENTTIME
4:                                      $\triangleright$  Replenish tokens
5:    $elapsed \leftarrow now - bucket.lastUpdate$ 
6:    $bucket.tokens \leftarrow \min(bucket.capacity,$ 
7:      $bucket.tokens + bucket.rate \times elapsed)$ 
8:    $bucket.lastUpdate \leftarrow now$ 
9:   if  $bucket.tokens \geq cost$  then
10:     $bucket.tokens \leftarrow bucket.tokens - cost$ 
11:    SAVEBUCKET( $key, bucket$ )
12:    return ALLOWED
13:  else
14:     $retryAfter \leftarrow (cost - bucket.tokens)/bucket.rate$ 
15:    return RATELIMITED( $retryAfter$ )
16:  end if
17: end function
```

4.2 Rate Limit Tiers

Table 1: Rate Limit Tiers

Tier	Rate (req/s)	Burst	Daily Limit
Free	10	20	10,000
Starter	50	100	100,000
Growth	200	500	1,000,000
Enterprise	1,000	2,000	Unlimited

4.3 Distributed State

Rate limit state is stored in Redis with the following structure:

Listing 6: Redis Rate Limit State

```
1 HSET ratelimit:{key} tokens {count} lastUpdate {timestamp}
2 EXPIRE ratelimit:{key} 3600
```

Atomic update using Lua script:

Listing 7: Atomic Rate Limit Check

```
1 local key = KEYS[1]
2 local rate = tonumber(ARGV[1])
3 local capacity = tonumber(ARGV[2])
4 local cost = tonumber(ARGV[3])
5 local now = tonumber(ARGV[4])
6
7 local bucket = redis.call('HMGET', key, 'tokens', 'lastUpdate')
8 local tokens = tonumber(bucket[1]) or capacity
9 local lastUpdate = tonumber(bucket[2]) or now
10
11 -- Replenish
12 local elapsed = now - lastUpdate
13 tokens = math.min(capacity, tokens + rate * elapsed)
14
15 if tokens >= cost then
16     tokens = tokens - cost
17     redis.call('HMSET', key, 'tokens', tokens, 'lastUpdate', now)
18     redis.call('EXPIRE', key, 3600)
19     return {1, tokens}
20 else
21     local retryAfter = (cost - tokens) / rate
22     return {0, retryAfter}
23 end
```

4.4 Rate Limit Headers

Responses include rate limit information:

Listing 8: Rate Limit Headers

```
1 X-RateLimit-Limit: 1000
2 X-RateLimit-Remaining: 742
3 X-RateLimit-Reset: 1548892800
4 Retry-After: 30 # Only on 429 responses
```

5 Caching Architecture

5.1 Cache Layers

1. **Edge Cache:** CDN caches public, immutable resources
2. **Gateway Cache:** Redis caches frequently accessed data
3. **Origin Cache:** Service-level caching for computed results

5.2 Cache Key Generation

Definition 5.1 (Cache Key). *For request req, the cache key is:*

$$key(req) = hash(req.method, req.path, sort(req.query), req.vary_headers) \quad (5)$$

5.3 Cache Invalidation

We implement a publish-subscribe invalidation protocol:

Algorithm 3 Cache Invalidation

```

1: function INVALIDATECACHE(resource)
2:   keys  $\leftarrow$  COMPUTEAFFECTEDKEYS(resource)
3:   for key  $\in$  keys do
4:     DELETEFROMCACHE(key)
5:   end for
6:   PUBLISH(invalidation_channel, resource)
7: end function
8: function ONRESOURCECHANGE(event)
9:   resource  $\leftarrow$  event.resource
10:  INVALIDATECACHE(resource)
11:  for related  $\in$  relationships(resource) do
12:    INVALIDATECACHE(related)
13:  end for
14: end function

```

5.4 Cache Coherence

Theorem 5.2 (Cache Consistency). *For any read r at time t following a write w at time $t_w < t$:*

$$t - t_w > \delta_{\text{invalidation}} \Rightarrow \text{read}(r) \text{ reflects write}(w) \quad (6)$$

where $\delta_{\text{invalidation}}$ is the maximum invalidation propagation delay.

In practice, $\delta_{\text{invalidation}} < 100\text{ms}$ with our pub-sub architecture.

5.5 Conditional Requests

Support for conditional requests reduces bandwidth:

Listing 9: Conditional Request

```

1 # Client sends:
2 GET /v1/products/123
3 If-None-Match: "abc123"
4
5 # Server responds:
6 HTTP/1.1 304 Not Modified
7 ETag: "abc123"

```

6 Error Handling

6.1 Error Response Format

Listing 10: Error Response

```

1 {
2   "error": {
3     "type": "invalid_request_error",
4     "code": "parameter_invalid",

```



```

5     "message": "The_price_must_be_a_positive_number",
6     "param": "price",
7     "doc_url": "https://docs.hanzo.ai/errors#parameter_invalid"
8 },
9 "request_id": "req_abc123"
10 }

```

6.2 Error Categories

Table 2: Error Types

Type	HTTP Status	Description
invalid_request_error	400	Malformed request
authentication_error	401	Invalid credentials
authorization_error	403	Insufficient permissions
not_found_error	404	Resource not found
rate_limit_error	429	Quota exceeded
api_error	500	Internal error

6.3 Idempotency

Write operations support idempotency keys:

Listing 11: Idempotent Request

```

1 POST /v1/orders
2 Idempotency-Key: order_123_attempt_1
3 Content-Type: application/json
4
5 {"items": [...], "customer": "cus_abc"}

```

Algorithm 4 Idempotent Request Processing

```

1: function PROCESSIDEMPOTENT(request)
2:   key  $\leftarrow$  request.headers[IdempotencyKey]
3:   if key  $\neq \perp$  then
4:     cached  $\leftarrow$  GETIDEMPOTENCYCACHE(key)
5:     if cached  $\neq \perp$  then
6:       if cached.requestHash = hash(request.body) then
7:         return cached.response
8:       else
9:         return IDEMPOTENCYKEYREUSED
10:      end if
11:    end if
12:  end if
13:  response  $\leftarrow$  PROCESSREQUEST(request)
14:  if key  $\neq \perp$  then
15:    SETIDEMPOTENCYCACHE(key, hash(request.body), response)
16:  end if
17:  return response
18: end function

```

7 Implementation

7.1 Technology Stack

- **Gateway:** Go with net/http
- **GraphQL:** gqlgen for schema-first development
- **Cache:** Redis Cluster
- **Load Balancing:** HAProxy with consistent hashing
- **Observability:** Prometheus, Jaeger, structured logging

7.2 Horizontal Scaling

Gateway instances are stateless and horizontally scalable:

$$throughput = n \times throughput_{single} \quad (7)$$

where n is the number of gateway instances.

8 Evaluation

8.1 Performance Benchmarks

Table 3: Gateway Performance

Metric	P50	P95	P99	Max
REST Latency	8ms	15ms	23ms	89ms
GraphQL Latency	12ms	28ms	45ms	156ms
Rate Limit Check	0.3ms	0.8ms	1.2ms	3ms
Cache Lookup	0.5ms	1.1ms	2.0ms	8ms

8.2 Throughput

Peak sustained throughput: 180,000 requests/second

Table 4: Throughput by Request Type

Request Type	Throughput (req/s)
REST GET (cached)	85,000
REST GET (uncached)	42,000
REST POST/PUT	28,000
GraphQL Query	18,000
GraphQL Mutation	7,000

8.3 Cache Performance

- Overall cache hit rate: 94%
- Product data hit rate: 97%
- Order data hit rate: 23% (low due to high write frequency)
- Average cache entry TTL: 5 minutes

8.4 Availability

Over 18-month production period:

- Uptime: 99.99%
- Total downtime: 52 minutes
- Incidents: 3 (all resolved within 20 minutes)

9 Related Work

API gateway design has evolved significantly. Kong [1] provides plugin-based extensibility. AWS API Gateway [2] offers serverless integration. GraphQL [3] introduced flexible querying, with Apollo Server [4] providing production implementations.

Rate limiting algorithms are well-studied. The token bucket [5] and leaky bucket [6] provide complementary approaches. Our distributed implementation draws from Redis-based patterns [7].

10 Conclusion

The Hanzo API Gateway provides a unified interface for e-commerce operations, supporting both REST and GraphQL with consistent authentication, rate limiting, and caching. The system achieves 180K requests/second throughput with sub-25ms P99 latency and 99.99% availability. Key innovations include GraphQL complexity analysis, distributed token bucket rate limiting, and a cache coherence protocol for commerce data.

Future directions include HTTP/3 support, edge computing for reduced latency, and machine learning-based anomaly detection for API abuse.

References

- [1] Kong Inc. Kong API Gateway. <https://konghq.com>, 2019.
- [2] Amazon Web Services. Amazon API Gateway. <https://aws.amazon.com/api-gateway>, 2019.
- [3] Facebook. GraphQL Specification. <https://graphql.org>, 2015.
- [4] Apollo GraphQL. Apollo Server. <https://apollographql.com>, 2019.
- [5] J. Turner. New directions in communications (or which way to the information age?). IEEE Communications Magazine, 24(10):8-15, 1986.
- [6] S. Bellovin. Leaky bucket counters and the low water mark. Computer Networks, 31(5):505-519, 1999.
- [7] Redis Labs. Rate limiting with Redis. Technical Documentation, 2019.