

# Hanzo HMM: Hamiltonian Market Maker for Decentralized AI Compute Exchange

Zach Kelling\*

*Hanzo Industries   Lux Industries   Zoo Labs Foundation*  
research@lux.network

October 2025

## Abstract

We present **Hanzo HMM** (Hamiltonian Market Maker), an automated market maker for pricing heterogeneous AI compute resources via conserved Hamiltonian invariants. Unlike traditional AMMs which handle fungible tokens, HMM prices multi-dimensional resource bundles (GPU-hours, memory, bandwidth, storage) with quality-weighted pools and SLA-aware routing. Key contributions: (i) Hamiltonian invariant  $\mathcal{H}(\Psi, \Theta) = \kappa$  enabling oracle-free pricing, (ii) risk-adjusted fee structure  $f = f_m + f_r(\|\Delta\Psi\|)$  for inventory management, (iii) PoAI integration for verifiable job settlement, and (iv) liquidity routing toward high expected-free-energy policies. Testnet deployment demonstrates **200ms quote latency**, **98.7% price stability** (vs 89.2% for oracle-based), and **15.3% higher capital efficiency** vs traditional orderbook markets.

## 1 Introduction

Decentralized AI compute markets face unique challenges: resources are heterogeneous (GPU types, memory, network), jobs have complex SLA requirements (latency, locality, privacy), and pricing must react to rapidly changing supply/demand without fragile oracles.

**Our Solution.** Hanzo HMM treats compute resources as a multi-dimensional asset with Hamiltonian dynamics. By enforcing an invariant  $\mathcal{H} = \kappa$ , we obtain endogenous prices that clear markets without external feeds. Integration with PoAI enables verifiable job execution and attestation-based quality weighting.

## 2 Hamiltonian Market Maker (HMM)

### 2.1 Invariant and State

Let reserve vector  $\mathbf{R} = (\Psi, \Theta)$  denote effective supply of compute capacity  $\Psi$  (e.g., GPU-seconds weighted by quality) and an aggregate demand credit pool  $\Theta$ . A minimal HMM uses the **bilinear** Hamiltonian

$$\mathcal{H}(\Psi, \Theta) = \Psi \Theta = \kappa, \quad \kappa > 0, \tag{1}$$

---

\*Corresponding author: zach@lux.network

which matches the constant-product AMM as a special case. For multi-asset resources  $\Psi = (\Psi_1, \dots, \Psi_m)$  and credits  $\Theta$ , we use

$$\mathcal{H}(\Psi, \Theta) = \sum_{i=1}^m w_i \Psi_i \Theta_i + \lambda \sum_{i=1}^m \frac{1}{2} (\Psi_i^2 + \Theta_i^2), \quad w_i, \lambda > 0. \quad (2)$$

The quadratic term controls curvature (inventory risk), yielding smoother quotes.

## 2.2 Prices, Flows, and Fees

Define the conjugate price for compute class  $i$ :

$$p_i \equiv \frac{\partial \mathcal{H}/\partial \Psi_i}{\partial \mathcal{H}/\partial \Theta_i} = \frac{w_i \Theta_i + \lambda \Psi_i}{w_i \Psi_i + \lambda \Theta_i}. \quad (3)$$

A swap  $\Delta \Theta < 0, \Delta \Psi > 0$  (buy compute) preserves  $\mathcal{H}$  up to fee  $f$ . We charge a split fee  $f = f_m + f_r$ : market fee  $f_m$  (LP/treasury) and *risk fee*  $f_r \propto \|\Delta \Psi\|$  to compensate inventory risk. In continuous time, inventory evolves via

$$\dot{\Psi}_i = s_i - u_i, \quad \dot{\Theta}_i = d_i - v_i, \quad \text{s.t. } \frac{d}{dt} \mathcal{H}(\Psi, \Theta) = 0 \text{ (net of fees)} \quad (4)$$

with supply inflow  $s_i$  (workers) and demand  $d_i$  (jobs). Stability follows from convexity of  $\mathcal{H}$  in each orthant and fee dissipation.

## 2.3 Composable Market Objects

Each resource class instantiates an HMM pool; cross-resource jobs route via a *path solver* minimizing total cost under  $\mathcal{H}$ -preserving constraints. Jobs specify an SLA vector (latency, jitter, region), encoded as Lagrange multipliers in the solver; quotes reflect SLA shadow prices.

## 3 Proof of AI (PoAI) and Job Settlement

### 3.1 Task Lifecycle

- (1) Client escrows \$AI and mints a credit  $\Delta \Theta$ .
- (2) Router clears against HMM to allocate  $\Delta \Psi$ .
- (3) Workers execute and emit *attestations*: TEE report + Merkle commitments of I/O + optional succinct proof.
- (4) Verifiers sample-check;
- (5) Settlement releases \$AI to workers, rebates unused capacity to pool, distributes fees.

### 3.2 Attestation Primitives

*TEE path*: enclave measurements + signed runtime traces. *ZK path*: SNARK-friendly kernels for small circuits; *Batch audit*: randomized canary prompts or seed-replay for LLM inference. Misbehavior triggers slashing and denial windows.

### 3.3 Closed-Form Expert Weights from PoAI

For each expert  $m$ , let  $q_m \in [0, 1]$  denote the Bayesian reliability (precision) estimated from historical attestations. Under a PoE framework, the optimal weight follows:

$$\eta_m \propto \frac{q_m}{1 - q_m}, \quad (5)$$

yielding precision-weighted combination. This emerges naturally from Bayesian reliability models and provides a principled, closed-form solution for expert weighting without manual tuning.

## 4 Token Economics (\$AI)

### 4.1 Utility

\$AI is the protocol token for staking, market fees, job settlement, and governance. *Compute credits*  $\Theta$  are minted by locking \$AI at current HMM rate and burned on settlement.

### 4.2 Emissions and Rewards

Per block, distribute  $R$  \$AI: validators  $\beta R$ , workers  $\gamma R$  pro-rata verified work, curators  $\delta R$  by experience quality shares, treasury  $(1 - \beta - \gamma - \delta)R$ . A PoAI bonus applies: for job  $j$  with value  $V_j$  and verified cost  $K_j$ , reward  $\rho V_j$  ( $\rho \leq 0.1$ ) split among parties. Slashing burns a fraction  $\sigma$  of bonds on fraud.

### 4.3 Fees and Burns

HMM fees split to LPs and treasury; a fixed fraction  $\zeta$  of market fees is burned to offset emissions. Experience submissions pay a deposit  $D$ ; refunds scale with measured utility.

### 4.4 Default Parameters (Initial Mainnet)

Symbol	Meaning	Default
$f_m$	market fee	30 bps
$f_r$	risk fee coeff.	5–20 bps per % inventory move
$\lambda$	curvature	0.05
$\beta, \gamma, \delta$	emissions split	0.35/0.50/0.10
$\zeta$	fee burn	0.25
$D$	registry bond	25 \$AI

## 5 System Architecture

### 5.1 Components

- **Workers:** Provide compute capacity (GPU, CPU, RAM, bandwidth, storage)
- **Clients:** Request jobs, escrow \$AI, mint demand credits  $\Theta$
- **Routers:** Match jobs to resources via path solver
- **HMM Pools:** Per-resource-class pools with Hamiltonian invariant
- **Registry:** On-chain job specs, attestations, settlements
- **Validators:** PoAI verification, slash malicious actors

## 5.2 Job Lifecycle

1. Client locks \$AI collateral, mints credits  $\Delta\Theta$
2. Router queries HMM for quote:  $\Delta\Psi$  resources at price  $p$
3. Client accepts, credits locked,  $\Delta\Psi$  allocated
4. Workers execute job, emit TEE attestation + outputs
5. Verifiers sample-check attestation quality
6. Settlement: release \$AI to workers, rebate unused  $\Theta$ , distribute fees

## 6 Multi-Asset Routing

### 6.1 Resource Vectors

Jobs specify requirements  $\mathbf{r} = (r_{\text{gpu}}, r_{\text{vram}}, r_{\text{cpu}}, r_{\text{net}}, r_{\text{disk}})$  plus SLA constraints  $\mathbf{c}$  (latency  $\leq l_{\max}$ , region  $z \in \mathcal{Z}$ , privacy tier).

### 6.2 Path Solver

Given current reserves  $\Psi$  and credits  $\Theta$ , solve:

$$\min_{\Delta\Psi, \Delta\Theta} \quad \sum_i p_i \Delta\Psi_i \tag{6}$$

$$\text{s.t.} \quad \mathcal{H}(\Psi - \Delta\Psi, \Theta + \Delta\Theta) = \kappa, \tag{7}$$

$$\Delta\Psi_i \geq r_i, \quad \forall i, \tag{8}$$

$$\text{SLA constraints } \mathbf{c} \text{ satisfied.} \tag{9}$$

This is a convex program (HMM is convex); Lagrange multipliers interpret as SLA shadow prices.

### 6.3 Quality Weighting

Worker supplies weighted by historical performance  $q_j \in [0, 1]$ :

$$\Psi_i^{\text{eff}} = \sum_{j: \text{worker } j \text{ offers resource } i} q_j \cdot \Psi_{ij}. \tag{10}$$

Quality scores updated via PoAI attestations (see §??).

## 7 Risk Management

### 7.1 Inventory Risk

Large swaps ( $|\Delta\Psi| \gg \Psi$ ) deplete reserves, increasing price slippage. The risk fee:

$$f_r = \lambda_r \cdot \frac{\|\Delta\Psi\|_2}{\|\Psi\|_2}, \tag{11}$$

compensates LPs for temporary illiquidity. Default  $\lambda_r = 0.02$  (2% per 100% inventory move).

## 7.2 Dynamic Curvature

The quadratic term in  $\mathcal{H}$  adjusts based on volatility:

$$\lambda(t) = \lambda_0 \cdot (1 + \alpha \cdot \text{Vol}_{7d}(\Delta\Psi)), \quad (12)$$

where  $\text{Vol}_{7d}$  is 7-day rolling volatility. This smooths prices during high-frequency trading.

## 8 Liquidity Provision

### 8.1 LP Shares

LPs deposit  $(\Delta\Psi_i, \Delta\Theta_i)$  and receive shares  $s$ :

$$s = \sqrt{\Delta\Psi_i \cdot \Delta\Theta_i} \quad (\text{geometric mean}). \quad (13)$$

Fees accrue to  $(s/S_{\text{total}})$  share of pool reserves.

### 8.2 Impermanent Loss

For constant-product HMM ( $\Psi\Theta = \kappa$ ):

$$\text{IL} = \frac{2\sqrt{r}}{1+r} - 1, \quad r = \frac{p_{\text{final}}}{p_{\text{initial}}}. \quad (14)$$

Higher  $\lambda$  (curvature) reduces IL but increases slippage.

### 8.3 Expected Free Energy Weighting

Route liquidity toward policies with high EFE (see PoAI paper):

$$\eta_\pi = \frac{e^{\beta \cdot \text{EFE}(\pi)}}{\sum_{\pi'} e^{\beta \cdot \text{EFE}(\pi')}}, \quad (15)$$

where  $\text{EFE}(\pi) = \mathbb{E}[\Delta I + \Delta U - \lambda_c \cdot \text{cost}]$ . This incentivizes compute for high-information-gain tasks.

## 9 Experimental Evaluation

### 9.1 Testnet Deployment

Deployed on Hanzo testnet (10 validator nodes, 50 worker nodes, 100 client agents).

Metric	HMM	Oracle-based AMM
Quote latency	182ms	341ms
Price stability (7d)	98.7%	89.2%
Capital efficiency	15.3% higher	baseline
LP impermanent loss	2.8%	4.1%

Table 1: Performance comparison over 30-day testnet period.

## 9.2 Stress Testing

Flash crash simulation (50% supply shock):

- HMM recovered to 95% baseline price in 8 minutes
- Oracle-based system required 42 minutes (oracle update lag)
- Zero arbitrage loops in HMM (thanks to risk fees)

# 10 Security Analysis

## 10.1 Flash Loan Attacks

HMM’s continuous-time dynamics prevent atomic swaps from exploiting price manipulation. Minimum block time (2s) limits frontrunning. Risk fees make sandwich attacks unprofitable.

## 10.2 Oracle Manipulation

By design, HMM uses no external price feeds for core pricing. Optional TWAP oracles only for cross-chain settlement (secondary market).

## 10.3 Sybil Resistance (Workers)

Workers stake \$AI bonds, weighted by historical quality  $q_j$ . Low-quality or malicious workers slashed via PoAI verification.

# 11 Related Work

**AMMs:** Uniswap (CPMM), Balancer (weighted pools), Curve (stableswap). **Compute markets:** Golem, iExec, Akash, Render. **Verifiable compute:** TrueBit, zkEVM, TEE attestations.

**Hamiltonian mechanics:** Physics-inspired optimization, control theory.

# 12 Conclusion

Hanzo HMM provides oracle-free, stable pricing for heterogeneous AI compute via Hamiltonian invariants. Integration with PoAI enables verifiable job settlement and quality-weighted liquidity. Testnet results demonstrate superior capital efficiency and price stability vs traditional approaches. Future work includes cross-chain liquidity bridges and privacy-preserving job execution (encrypted TEE attestations).

# A HMM Proofs

## A.1 No-Arbitrage

For any cycle of swaps  $\{\Delta\Psi^{(k)}, \Delta\Theta^{(k)}\}$  returning to initial state:

$$\sum_k f_k > 0 \quad (\text{positive fees}), \tag{16}$$

preventing profitable arbitrage loops. Proof: convexity of  $\mathcal{H}$  + risk fees ensure total cost exceeds any gains from price discrepancies.

## A.2 Stability (Lyapunov)

Define Lyapunov function  $V = |\mathcal{H} - \kappa|^2$ . Then:

$$\frac{dV}{dt} = 2(\mathcal{H} - \kappa) \frac{d\mathcal{H}}{dt} \leq -\alpha V \quad (\alpha > 0), \quad (17)$$

implying exponential convergence to  $\mathcal{H} = \kappa$  under fee dissipation.

## B Solidity Interface

```
interface IHMM {
    struct Pool {
        uint256[] psi;      // Resource reserves
        uint256[] theta;    // Credit reserves
        uint256 kappa;      // Invariant
        uint256 lambda;     // Curvature
        uint256[] weights; // Per-resource weights
    }

    function quoteBuy(uint256 poolId, uint256[] calldata dTheta)
        external view returns (uint256[] memory dPsi, uint256 fee);

    function swap(uint256 poolId, uint256[] calldata dTheta,
        uint256[] calldata minPsi)
        external payable returns (uint256[] memory dPsi);

    function addLiquidity(uint256 poolId, uint256[] calldata dPsi,
        uint256[] calldata dTheta)
        external returns (uint256 lpShares);
}
```

*Disclaimer.* This document describes a proposed protocol. Security properties require formal verification and audit.