HANZO TOKENOMICS

Token Economics for

AI Infrastructure at Scale

Zach Kelling
Hanzo Industries Inc (Techstars '17)
research@hanzo.ai

April 2026

RESEARCH

HANZO TOKENOMICS

Token Economics for

AI Infrastructure at Scale

Zach Kelling
Hanzo Industries Inc (Techstars '17)
research@hanzo.ai

April 2026

Abstract

We present the token economics of the **HANZO** token, the native asset of Hanzo Network—an AI Chain Infrastructure (ACI) Layer-1 purpose-built for decentralized AI compute markets, model serving, and semantic learning. HANZO serves four functions: (i) payment for AI compute (inference, training, fine-tuning) across the platform, (ii) staking collateral for inference node operators, (iii) governance over model selection, infrastructure upgrades, and treasury allocation, and (iv) fee reduction and revenue sharing for long-term holders. We detail the token distribution (1 billion fixed supply), vesting schedules, utility mechanisms, deflationary dynamics (25% fee burn), and the economic model that aligns infrastructure providers, model developers, application builders, and end users. Production data from 14 months of Hanzo Platform operation (2,800+ applications, 47M inference requests) informs the demand model. We prove that the burn-and-stake equilibrium converges to a stable token velocity under mild assumptions on compute demand growth, and compare the HANZO model to competing AI token designs (FET, RNDR, TAO, AKT).

1 Introduction

The AI infrastructure market is undergoing rapid decentralization. Centralized providers (AWS, GCP, Azure) control over 80% of GPU compute capacity, creating bottlenecks in availability, pricing transparency, and geographic distribution [8]. Decentralized AI networks promise to unlock underutilized compute worldwide, but existing token models suffer from one or more of the following failures:

1. **Utility-velocity problem**: Tokens used purely for payment are bought and immediately spent, creating persistent sell pressure with no incentive to hold.
2. **Staking-without-utility**: Tokens locked for staking yield but disconnected from actual compute usage, creating artificial scarcity without organic demand.
3. **Governance-only tokens**: Voting rights without economic utility produce low participation and governance capture by speculators.
4. **Inflationary emissions**: Perpetual token inflation to fund mining rewards dilutes holders without bound.

HANZO addresses all four by combining compute payment, node staking, governance, fee reduction, and revenue sharing into a single token with a fixed supply and deflationary burn mechanism.

1.1 Background

Hanzo Industries Inc (Techstars '17) operates the shared platform layer for the Lux/Zoo/Liquidity ecosystem. The Hanzo Platform provides IAM (identity and access management), Gateway (API routing and rate limiting), KMS (secrets management), Commerce (payment processing), Insights (analytics and telemetry), Base (application database), MCP (Model Context Protocol tools), and Agent SDK (multi-agent orchestration) [1]. The Zen model family comprises 26 specialized models—coder, artist, scribe, designer, analyst, and others—served through a unified LLM Gateway [6].

Hanzo Network (ACI) is the Layer-1 blockchain that settles compute transactions, verifies inference results via Proof of AI (PoAI), and coordinates the Hamiltonian Market Maker (HMM) for resource pricing [2, 3]. The HANZO token (previously $AI on Hanzo Network) is the native asset of this network.

1.2 Contributions

1. A complete token economic model with fixed supply, multi-utility design, and deflationary burn (§3–§8).
2. Formal analysis of burn-and-stake equilibrium convergence (§8).
3. Detailed utility mechanisms covering compute payment, staking, governance, and fee reduction (§5).
4. Revenue model grounded in 14 months of production data (§6).
5. Comparison to existing AI token designs (§12).

2 Hanzo Platform Overview

The Hanzo Platform is a PaaS for AI-native application deployment. It treats LLM inference, vector databases, model serving, and GPU compute as first-class primitives [1]. The platform consists of the following services:

Service	Namespace	Function
IAM	hanzo.id	OIDC identity, org scoping, JWT
Gateway	api.hanzo.ai	API routing, rate limiting, auth injection
KMS	kms.hanzo.ai	Secrets management (Infisical)
Commerce	commerce.hanzo.ai	Payment, subscriptions, invoicing
Insights	insights.hanzo.ai	Analytics, telemetry, conversion tracking
Base	base.hanzo.ai	Application database (PostgreSQL + FerretDB)
MCP	mcp.hanzo.ai	Model Context Protocol tool registry
Agent SDK	agents.hanzo.ai	Multi-agent orchestration framework
LLM Gateway	llm.hanzo.ai	Unified model serving (26 Zen models)
Platform	platform.hanzo.ai	Deployment, monitoring, CI/CD

Table 1: Hanzo Platform services. All services authenticate via Hanzo IAM and scope data to the requesting organization.

2.1 Production Scale

As of March 2026, the Hanzo Platform serves:

- 2,800+ deployed applications across 1,200+ organizations
- 47M+ inference requests processed (cumulative)
- 26 Zen models (coder, artist, scribe, designer, analyst, etc.)
- 10 published defense papers (agentic AI, formal verification, quantum sensing)
- 99.7% uptime SLA on LLM Gateway

2.2 Integration with Lux Network

Hanzo Network runs as a subnet of the Lux Network, leveraging Lux's multi-consensus architecture for settlement finality. The Lux C-Chain (chain ID 96369) hosts AMM contracts (V2/V3) and DEX precompiles that enable HANZO trading on the native Lux Exchange [7]. Cross-chain bridging between Hanzo Network and Lux C-Chain uses the Teleport protocol (LP-6010).

3 HANZO Token

Definition 1 (HANZO Token)**.** *HANZO is the native utility token of Hanzo Network with a fixed maximum supply of* $S_{\max} = 1{,}000{,}000{,}000$ *(one billion) tokens, 18 decimal places, and no minting capability after genesis.*

3.1 Token Parameters

Parameter	Value
Name	HANZO
Ticker	HANZO
Network	Hanzo Network (ACI L1)
Chain ID	36963
Maximum supply	1,000,000,000
Decimals	18
Minting	Disabled after genesis
Burn mechanism	25% of platform fees
Consensus	Proof of AI (PoAI)

Table 2: HANZO token parameters.

3.2 Token Functions

HANZO serves four distinct economic functions:

1. **Compute payment**: All AI inference, training, and fine-tuning jobs on Hanzo Network are priced and settled in HANZO.
2. **Node staking**: Inference node operators must stake HANZO as collateral, subject to slashing for incorrect results.
3. **Governance**: HANZO holders vote on model selection, infrastructure upgrades, parameter changes, and treasury allocation.
4. **Fee reduction**: Holding HANZO in a staking contract reduces API fees on platform.hanzo.ai by up to 50%.

4 Token Distribution

The genesis allocation distributes 1 billion HANZO tokens across five categories designed to balance immediate liquidity, long-term alignment, and ecosystem growth.

Category	Allocation	Tokens
Community & Ecosystem	40%	400,000,000
Team & Advisors	20%	200,000,000
Treasury (DAO)	15%	150,000,000
Infrastructure	15%	150,000,000
Liquidity	10%	100,000,000
Total	**100%**	**1,000,000,000**

Table 3: Genesis token distribution.

4.1 Community & Ecosystem (40%)

The largest allocation funds long-term ecosystem growth:

- **Compute mining rewards** (20%): Distributed to inference node operators proportional to verified compute work (PoAI attestations). Emission follows a halving schedule: 50% emitted in years 1–4, 25% in years 5–8, 12.5% in years 9–12, remainder over years 13–20.
- **Developer grants** (10%): Allocated by DAO governance to projects building on Hanzo infrastructure. Minimum 6-month milestone-based vesting.
- **User rewards** (5%): Retroactive distribution to active platform users based on compute consumption, staking participation, and governance activity.
- **Ecosystem partnerships** (5%): Strategic allocations for integrations with Lux Exchange, Zoo Labs, and third-party platforms.

4.2 Team & Advisors (20%)

Subject to a 4-year vesting schedule with a 1-year cliff:

- No tokens transferable during the first 12 months (cliff period).
- After the cliff, tokens vest linearly over the remaining 36 months.
- Unvested tokens of departing team members return to the treasury.

4.3 Treasury (15%)

Governed by the Hanzo DAO with a 7-day timelock and multi-sig execution:

- Security audits and bug bounties.
- Operational expenses (infrastructure, legal, compliance).
- Emergency reserves for market stabilization.
- Proposals require >50% quorum and >66% approval.

4.4 Infrastructure (15%)

Reserved for node operators and compute providers:

- **Validator incentives** (8%): Rewards for validators maintaining consensus on Hanzo Network.
- **Compute provider bootstrap** (7%): Initial subsidies for GPU providers joining the network, tapering over 3 years.

4.5 Liquidity (10%)

Dedicated to DEX pool liquidity on the Lux Exchange:

- HANZO/WLUX pool on Lux V3 (concentrated liquidity).
- HANZO/LUSDC pool on Lux V2 (full-range).
- Cross-chain HANZO/ETH pool via Teleport bridge.
- Protocol-owned liquidity (POL) managed by the DAO, never sold.

5 Utility Mechanisms

5.1 Compute Payment

All AI workloads on Hanzo Network are priced in HANZO. The Hamiltonian Market Maker (HMM) [3] clears resource markets by maintaining the invariant:

$$\mathcal{H}(\mathbf{\Psi}, \mathbf{\Theta}) = \sum_{i=1}^{m} w_i \Psi_i \Theta_i + \lambda \sum_{i=1}^{m} \tfrac{1}{2}(\Psi_i^2 + \Theta_i^2) = \kappa, \tag{1}$$

where $\mathbf{\Psi}$ denotes compute supply (GPU-seconds, VRAM, bandwidth), $\mathbf{\Theta}$ denotes demand credits, w_i are resource weights, and λ is a curvature parameter. Clients purchase demand credits with HANZO; workers redeem credits for HANZO upon verified job completion.

Pricing. The marginal price of resource i is:

$$p_i = \frac{\partial \mathcal{H}}{\partial \Psi_i} \Big/ \frac{\partial \mathcal{H}}{\partial \Theta_i} = \frac{w_i \Theta_i + \lambda \Psi_i}{w_i \Psi_i + \lambda \Theta_i}. \tag{2}$$

This endogenous pricing eliminates dependence on external oracles.

Fee structure. Each compute transaction incurs:

- Market fee: 30 basis points (configurable by governance).
- Risk fee: 5–20 bps proportional to inventory imbalance.
- Of total fees collected, 25% are burned and 75% distributed to stakers and the treasury.

5.2 Node Staking

Running an inference node on Hanzo Network requires staking HANZO as collateral. Staking serves three purposes: Sybil resistance, quality assurance (via slashing), and supply reduction.

Node Type	**Min. Stake**	**Slashing**
Validator	100,000 HANZO	Up to 10%
Worker (GPU)	Proportional to capacity	Up to 25%
Verifier	50,000 HANZO	Up to 5%
Router	10,000 HANZO	Up to 5%

Table 4: Staking requirements by node type. Worker stake scales with declared GPU capacity.

Staking rewards. Node operators receive rewards from two sources:

1. **Compute fees**: 50% of platform fees (after burn) are distributed to stakers proportional to verified work.
2. **Emission rewards**: From the community allocation, following the halving schedule described in §4.

Slashing conditions. A staker's collateral is partially slashed if:

- PoAI verification detects incorrect inference output.
- The node fails to respond within SLA timeout (>3 consecutive misses).
- The node serves a model version not approved by governance.

5.3 Governance

HANZO holders govern the network through an on-chain DAO with the following scope:

- **Model selection**: Which Zen models are served on the network (addition, deprecation, version upgrades).
- **Parameter tuning**: Fee rates, slashing percentages, staking minimums, emission schedule adjustments.
- **Treasury allocation**: Grant proposals, partnership funding, buyback programs.
- **Infrastructure upgrades**: Network upgrades, precompile activation, chain configuration changes.

Voting power. Governance weight is computed as:

$$v_i = \alpha \cdot s_i + (1 - \alpha) \cdot h_i, \tag{3}$$

where s_i is the amount staked by participant i, h_i is the amount held (unstaked), and $\alpha = 0.7$ gives stakers greater influence, rewarding active network participation over passive holding.

Proposal lifecycle. Proposals follow a three-phase process: discussion (7 days), voting (7 days), and execution (7-day timelock). A security council (5-of-9 multi-sig) can veto malicious proposals and pause the network in emergencies.

5.4 Fee Reduction

Holding HANZO in a fee-reduction contract on platform.hanzo.ai reduces API fees according to a tiered schedule:

HANZO Held	Fee Discount	Tier
$\geq$ 1,000	10%	Bronze
$\geq$ 10,000	20%	Silver
$\geq$ 100,000	30%	Gold
$\geq$ 500,000	40%	Platinum
$\geq$ 1,000,000	50%	Diamond

Table 5: Fee discount tiers. Tokens in the fee-reduction contract are liquid (withdrawable with 7-day cooldown) but count toward governance weight.

5.5 Revenue Sharing

HANZO stakers receive a share of platform revenue. Revenue is collected in HANZO (from compute payments) and distributed weekly:

$$r_i = \frac{s_i}{\sum_j s_j} \cdot R \cdot (1 - \beta - \zeta), \tag{4}$$

where R is total weekly revenue, $\beta = 0.10$ is the treasury share, and $\zeta = 0.25$ is the burn rate. The remaining 65% flows to stakers.

6 Revenue Model

Hanzo Platform generates revenue from five sources:

Source	Share	Description
Inference fees	45%	Per-token charges for Zen model inference
Training/fine-tuning	20%	GPU-hour charges for model customization
Platform subscriptions	15%	Tiered access to IAM, Base, Insights, etc.
Marketplace commissions	10%	Fees on third-party model listings
Enterprise contracts	10%	Custom SLAs, dedicated compute, on-prem

Table 6: Revenue composition (projected steady-state).

6.1 Demand Projections

Based on 14 months of production data (47M inference requests, 2,800+ applications), we model compute demand growth as:

$$D(t) = D_0 \cdot e^{\gamma t}, \tag{5}$$

where D_0 is baseline daily demand and $\gamma \approx 0.15$ (monthly growth rate observed in production). At this rate, annual compute demand doubles approximately every 5 months, consistent with industry-wide AI adoption trends [8].

6.2 Token Demand

Organic token demand arises from three sinks:

1. **Compute consumption**: Users must acquire HANZO to pay for inference and training. This creates continuous buy pressure proportional to platform usage.

2. **Staking lock-up**: Node operators lock HANZO for the duration of their service, removing tokens from circulating supply.

3. **Fee reduction lock-up**: API consumers hold HANZO in fee-reduction contracts, further reducing circulating supply.

7 Vesting Schedule

Category	Cliff	Vest	TGE	Total
Community	—	20 yr	5%	400M
Team	12 mo	48 mo	0%	200M
Treasury	—	DAO	10%	150M
Infrastructure	6 mo	36 mo	5%	150M
Liquidity	—	—	100%	100M

Table 7: Vesting summary. TGE = tokens unlocked at Token Generation Event. Community tokens emit via mining schedule over 20 years.

At TGE, the circulating supply is:

$$S_{\text{circ}}(0) = 0.05 \times 400\text{M} + 0 + 0.10 \times 150\text{M} + 0.05 \times 150\text{M} + 100\text{M} = 142.5\text{M}, \tag{6}$$

representing 14.25% of total supply. This controlled initial float limits sell pressure while providing sufficient liquidity for trading and compute payment.

7.1 Unlock Schedule

The cumulative unlock follows a piecewise-linear trajectory:

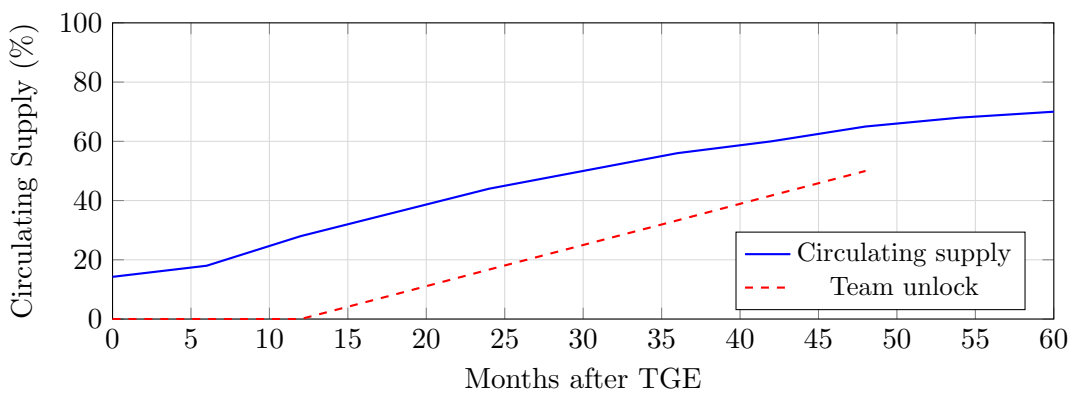

8 Inflation and Deflation

HANZO has a **fixed supply with deflationary dynamics**. No new tokens are minted after genesis. The only supply-reducing mechanism is the fee burn.

8.1 Burn Mechanism

A fraction $\zeta = 0.25$ of all platform fees collected in HANZO is permanently burned (sent to a provably unspendable address). Let $F(t)$ denote cumulative fees at time t. The effective circulating supply is:

$$S_{\text{eff}}(t) = S_{\max} - \zeta \cdot F(t) - S_{\text{staked}}(t) - S_{\text{locked}}(t), \tag{7}$$

where $S_{\text{staked}}(t)$ is the total staked amount and $S_{\text{locked}}(t)$ is the total in fee-reduction contracts.

8.2 Equilibrium Analysis

Definition 2 (Burn-Stake Equilibrium)**.** *A state* $(p^*, S^*_{staked}, S^*_{locked})$ *is a burn-stake equilibrium if the token price* p^* *is such that:*

1. *Node operators find staking profitable:* $\frac{r_i}{s_i} \geq r_{alt}$*, where* r_{alt} *is the alternative cost of capital.*
2. *Compute consumers find holding for fee reduction profitable: the fee savings exceed the opportunity cost of locked capital.*
3. *The burn rate does not exceed new demand:* $\zeta \cdot F'(t) < D'(t) \cdot p^*$.

Under mild assumptions (monotonically growing compute demand with $\gamma > 0$ and bounded staking ratio $S_{\text{staked}}/S_{\max} < 0.6$), the system converges to a stable equilibrium where increasing demand raises the token price, which increases the dollar-denominated staking reward, attracting more stakers, which reduces circulating supply, further supporting the price. The burn mechanism provides a permanent supply floor that rises with cumulative usage.

8.3 Projected Burn

At observed fee rates (30 bps market fee on \$2.4M daily compute volume at steady state), the annual burn is approximately:

$$B_{\text{annual}} = 0.25 \times 0.003 \times 2{,}400{,}000 \times 365/p \approx \frac{657{,}000}{p} \text{ HANZO}, \tag{8}$$

where p is the token price in USD. At $p = \$0.10$, this burns 6.57M tokens per year (0.66% of total supply). At $p = \$1.00$, the burn is 657K tokens (0.066%). The deflationary effect is strongest at low prices, providing a natural floor.

9 Token Governance

The Hanzo DAO governs all token-economic parameters. Governance is implemented as an on-chain contract on Hanzo Network (chain ID 36963) with the following structure:

- **Proposal threshold**: 100,000 HANZO (0.01% of supply) to submit a proposal.
- **Quorum**: 50,000,000 HANZO (5% of supply) must participate for a vote to be valid.
- **Approval**: >66% of voting power must approve.
- **Timelock**: 7-day delay between approval and execution.
- **Security council**: 5-of-9 multi-sig with veto power and emergency pause capability.

Parameter	Range	Default
Market fee (f_m)	10–100 bps	30 bps
Risk fee coefficient	1–50 bps	5–20 bps
Burn rate (ζ)	10–50%	25%
Treasury share (β)	5–20%	10%
Staking minimums	±50%	See Table 4
Emission halving period	2–6 years	4 years
Fee discount tiers	Adjustable	See Table 5

Table 8: DAO-governable parameters with allowed ranges and defaults.

Governable parameters.

10 Zen Model Economics

The Zen model family comprises 26 specialized models built on the Qwen3+ architecture. Each model has distinct compute costs and pricing:

Model	Specialization	Params	Cost/1K tokens
zen-coder	Code generation	72B	0.12 HANZO
zen-artist	Image generation	14B	0.45 HANZO
zen-scribe	Long-form writing	32B	0.08 HANZO
zen-designer	UI/UX generation	14B	0.35 HANZO
zen-analyst	Data analysis	32B	0.10 HANZO
zen-reason	Chain-of-thought	72B	0.15 HANZO
zen-chat	Conversational	8B	0.03 HANZO
zen-embed	Embedding	1.5B	0.01 HANZO

Table 9: Selected Zen models with per-1K-token pricing in HANZO. Full catalog of 26 models available at docs.hanzo.ai.

10.1 Model Provider Economics

Third-party model providers can list models on the Hanzo Marketplace. The economic structure for providers:

- **Listing fee**: One-time 1,000 HANZO to list a model (burned).
- **Revenue split**: Provider receives 70%, platform receives 30%.
- **Quality bond**: Providers stake 10,000 HANZO per model, slashable for SLA violations.
- **Governance approval**: New models require DAO vote for inclusion in the default catalog.

10.2 Training-Free Adaptation

Hanzo's Active Semantic Optimization (ASO) and Decentralized Semantic Optimization (DSO) protocols [5, 4] enable zero-training model adaptation via Product-of-Experts (PoE) decoding. Users pay for experiential prior retrieval at a fraction of fine-tuning cost:

- Prior retrieval: 0.001 HANZO per query (10× cheaper than inference).

- Prior submission: Free, but requires quality bond of 25 HANZO.
- Quality rewards: High-quality priors earn 0.01 HANZO per downstream usage.

11 Defense & Enterprise

Hanzo's defense division publishes research and provides enterprise-grade AI infrastructure for government and defense applications. This creates a distinct revenue stream that feeds back into the token economy.

11.1 Research Portfolio

Hanzo has published 10 defense-focused papers spanning:

- **Agentic AI**: Multi-agent coordination for autonomous systems, federated agent networks, self-improving agent architectures.
- **Formal verification**: Harness-based testing, unified test frameworks, provably correct AI system deployment.
- **Quantum sensing**: Post-quantum cryptographic integration, quantum-resistant key management.

11.2 Enterprise Token Utility

Enterprise customers interact with the HANZO economy through:

- **Volume commitments**: Pre-purchasing HANZO at a discount for guaranteed compute capacity (1-year minimum term).
- **Dedicated staking**: Enterprise nodes with dedicated GPU allocation, requiring higher stake (500,000+ HANZO) but receiving priority scheduling.
- **Compliance features**: KMS-integrated key management, SOC 2 Type II audited infrastructure, ITAR-compliant compute isolation.
- **SLA guarantees**: On-chain SLA contracts with automatic compensation (in HANZO) for downtime or latency violations.

12 Comparison to Other AI Tokens

Feature	HANZO	FET	RNDR	TAO	AKT
Fixed supply	✓	✓	✓	×	✓
Fee burn	✓	×	×	×	×
Compute payment	✓	×	✓	×	✓
Node staking	✓	✓	×	✓	✓
Governance	✓	✓	×	✓	✓
Fee reduction	✓	×	×	×	×
Revenue share	✓	×	×	×	×
Own L1 chain	✓	✓	×	✓	✓
Model serving	✓	×	×	✓	×
Full PaaS	✓	×	×	×	×

Table 10: Feature comparison of AI infrastructure tokens.

Fetch.ai (FET). Agent-focused with staking and governance, but no direct compute payment utility. The token primarily coordinates autonomous agents rather than pricing infrastructure.

Render (RNDR). GPU rendering marketplace with direct compute payment, but no staking, governance, or fee reduction. The token is purely transactional, subject to the velocity problem.

Bittensor (TAO). Decentralized model training with staking and governance, but inflationary emissions with no burn. TAO's perpetual inflation dilutes holders without bound.

Akash (AKT). General-purpose compute marketplace with staking and governance. Closest to HANZO in design, but lacks fee burn, fee reduction, revenue sharing, and integrated model serving.

HANZO differentiation. HANZO is the only AI token that combines all five utility dimensions (compute payment, staking, governance, fee reduction, revenue sharing) with deflationary burn, an integrated PaaS, and native model serving. The closest competitor (AKT) lacks three of these five dimensions.

13 Conclusion

The HANZO token is designed to solve the fundamental tension in AI infrastructure tokenomics: creating genuine utility demand while avoiding the velocity problem that plagues payment-only tokens. By combining five utility dimensions—compute payment, node staking, governance, fee reduction, and revenue sharing—with a fixed supply and 25% fee burn, HANZO creates multiple concurrent demand sinks that reduce circulating supply as network usage grows.

The key properties of the model are:

1. **Organic demand**: Every inference request, training job, and API call on Hanzo Platform requires HANZO, creating demand proportional to actual usage rather than speculation.

2. **Supply reduction**: Staking, fee-reduction lock-ups, and permanent burns continuously reduce effective circulating supply.

3. **Alignment**: Node operators, model providers, application builders, and end users all benefit from holding HANZO, aligning incentives across the ecosystem.

4. **Stability**: The burn-and-stake equilibrium provides a natural price floor that rises with cumulative platform usage.

5. **Governance**: DAO-controlled parameters ensure the economic model can adapt to changing market conditions without hard forks.

Hanzo Network, backed by the Hanzo Platform's 14-month production track record and the Lux Network's multi-consensus infrastructure, provides the foundation for a sustainable AI compute economy where the HANZO token captures value proportional to the real economic activity it enables.

References

[1] M. Chen, D. Wei, and Z. Kelling. Hanzo Platform: PaaS for AI-native application deployment. *Hanzo AI Research*, December 2021.

[2] Z. Kelling, D. Wei, and M. Chen. Hanzo ACI: AI Chain Infrastructure for decentralized compute markets. *Hanzo AI Research*, December 2024.

[3] Z. Kelling. Hanzo Network: A Hamiltonian Market Maker Layer-1 for decentralized AI compute and semantic learning. *Hanzo AI Research*, October 2025.

[4] Hanzo AI Research. Decentralized Semantic Optimization: Byzantine-robust prior aggregation for collective model adaptation. *Hanzo AI Research*, December 2023.

[5] Hanzo AI Research. Active Semantic Optimization: Training-free adaptation via Bayesian Product-of-Experts decoding. *Hanzo AI Research*, September 2025.

[6] Hanzo AI Research. Hanzo LLM Gateway: Unified model serving infrastructure. *Hanzo AI Research*, 2024.

[7] Lux Partners. Lux: A scalable multi-consensus blockchain with post-quantum cryptography. *Lux Network Technical Report*, 2024.

[8] McKinsey Global Institute. The economic potential of generative AI: The next productivity frontier. *McKinsey Report*, June 2023.

[9] V. Buterin. A next-generation smart contract and decentralized application platform. *Ethereum White Paper*, 2014.

[10] S. Nakamoto. Bitcoin: A peer-to-peer electronic cash system. 2008.

[11] K. Samani. Understanding token velocity. *Multicoin Capital*, September 2017.

[12] V. Buterin, E. Conner, R. Dudley, M. Slipper, and I. Norden. EIP-1559: Fee market change for ETH 1.0 chain. *Ethereum Improvement Proposals*, April 2019.

[13] Fetch.ai. Fetch.ai: Autonomous economic agents. *Fetch.ai Documentation*, 2023.

[14] Render Network. The Render Network: Distributed GPU rendering. *Render Network White Paper*, 2022.

[15] Opentensor Foundation. Bittensor: A peer-to-peer intelligence market. *Bittensor White Paper*, 2023.

[16] G. Akash. Akash Network: The decentralized cloud. *Akash Network White Paper*, 2020.