

# Zen AI Model Family

## Zen-Nano

Mobile/IoT Intelligence

Technical Whitepaper v1.0

Hanzo AI Research Team  
[research@hanzo.ai](mailto:research@hanzo.ai)

Zoo Labs Foundation  
[foundation@zoolabs.org](mailto:foundation@zoolabs.org)

September 2025

### Abstract

We present **Zen-Nano**, a 0.6B parameter model optimized for mobile/iot intelligence. Built upon zen-0.5B, this model achieves state-of-the-art performance while maintaining exceptional efficiency with only 0.6B active parameters. Supporting 64K thinking tokens for advanced reasoning, the model represents a significant advancement in democratizing AI through sustainable and efficient architectures.

## Contents

|          |                                    |          |
|----------|------------------------------------|----------|
| <b>1</b> | <b>Introduction</b>                | <b>2</b> |
| 1.1      | Key Innovations . . . . .          | 2        |
| <b>2</b> | <b>Architecture</b>                | <b>2</b> |
| 2.1      | Model Design . . . . .             | 2        |
| 2.2      | Technical Innovations . . . . .    | 2        |
| 2.2.1    | Mixture of Experts (MoE) . . . . . | 2        |
| 2.2.2    | Attention Mechanism . . . . .      | 2        |
| 2.2.3    | Thinking Mode . . . . .            | 2        |
| <b>3</b> | <b>Performance Benchmarks</b>      | <b>3</b> |
| 3.1      | Evaluation Results . . . . .       | 3        |
| 3.2      | Efficiency Metrics . . . . .       | 3        |
| <b>4</b> | <b>Training Methodology</b>        | <b>3</b> |
| 4.1      | Dataset . . . . .                  | 3        |
| 4.2      | Training Process . . . . .         | 3        |
| <b>5</b> | <b>Use Cases and Applications</b>  | <b>3</b> |
| 5.1      | Primary Applications . . . . .     | 3        |
| 5.2      | Integration Examples . . . . .     | 4        |
| <b>6</b> | <b>Environmental Impact</b>        | <b>4</b> |
| 6.1      | Sustainability Metrics . . . . .   | 4        |
| 6.2      | Green AI Commitment . . . . .      | 4        |

|                                      |          |
|--------------------------------------|----------|
| <b>7 Safety and Alignment</b>        | <b>4</b> |
| 7.1 Safety Measures . . . . .        | 4        |
| 7.2 Ethical Considerations . . . . . | 4        |
| <b>8 Deployment Options</b>          | <b>5</b> |
| 8.1 Available Formats . . . . .      | 5        |
| 8.2 Hardware Requirements . . . . .  | 5        |
| <b>9 Future Work</b>                 | <b>5</b> |
| 9.1 Planned Improvements . . . . .   | 5        |
| 9.2 Research Directions . . . . .    | 5        |
| <b>10 Conclusion</b>                 | <b>5</b> |
| <b>A Model Card</b>                  | <b>6</b> |

# 1 Introduction

The rapid advancement of artificial intelligence has created an unprecedented demand for models that balance capability with efficiency. **Zen-Nano** addresses this challenge by delivering enterprise-grade performance while maintaining a minimal computational footprint.

## 1.1 Key Innovations

- **Efficient Architecture:** 0.6B active parameters from 0.6B total
- **Specialized Training:** Optimized for mobile/iot intelligence
- **Extended Context:** 32K context window
- **Thinking Mode:** 64K thinking tokens

# 2 Architecture

## 2.1 Model Design

Zen-Nano is based on the zen-0.5B architecture with several key modifications:

| Component         | Specification |
|-------------------|---------------|
| Total Parameters  | 0.6B          |
| Active Parameters | 0.6B          |
| Base Model        | zen-0.5B      |
| Context Length    | 32K           |
| Thinking Tokens   | 64K           |
| Architecture Type | Transformer   |

Table 1: Zen-Nano Architecture Specifications

## 2.2 Technical Innovations

### 2.2.1 Mixture of Experts (MoE)

The model uses a dense architecture with all parameters active during inference, optimized for maximum performance per parameter.

### 2.2.2 Attention Mechanism

Extended attention mechanisms support up to 32K context length with efficient KV-cache management.

### 2.2.3 Thinking Mode

Advanced reasoning through extended thinking tokens (up to 64K), enabling:

- Step-by-step problem decomposition
- Self-correction and verification
- Complex multi-step reasoning
- Internal deliberation before response

## 3 Performance Benchmarks

### 3.1 Evaluation Results

| Benchmark | Score |
|-----------|-------|
| MMLU      | 51.7% |
| HumanEval | 22.6% |
| GSM8K     | 62.0% |
| HellaSwag | 59.5% |

Table 2: Language Understanding Benchmarks

### 3.2 Efficiency Metrics

| Metric                | Value          |
|-----------------------|----------------|
| Inference Speed       | 450 tokens/sec |
| Memory Usage (INT4)   | 2 GB           |
| Energy Efficiency     | 98% reduction  |
| Latency (First Token) | 15 ms          |

Table 3: Efficiency Metrics

## 4 Training Methodology

### 4.1 Dataset

The model was trained on a carefully curated dataset comprising:

- High-quality filtered web data (0.5TB)
- Domain-specific corpora for mobile/iot intelligence
- Synthetic data generation for edge cases
- Human feedback through RLHF

### 4.2 Training Process

1. **Pretraining:** 2 trillion tokens over 14 days on 8x A100
2. **Supervised Fine-tuning:** Task-specific optimization
3. **RLHF:** Alignment with human preferences
4. **Constitutional AI:** Safety and helpfulness optimization

## 5 Use Cases and Applications

### 5.1 Primary Applications

Conversational AI and chatbots

Content generation and summarization

Code completion and review

Educational assistance

Research and analysis

## 5.2 Integration Examples

```
1 from transformers import AutoModelForCausalLM, AutoTokenizer
2
3 # Load model and tokenizer
4 model = AutoModelForCausalLM.from_pretrained("zenlm/zen-nano-0.6b-
5   instruct")
6 tokenizer = AutoTokenizer.from_pretrained("zenlm/zen-nano-0.6b-instruct
7   ")
8
9 # Generate response
10 inputs = tokenizer("Explain quantum computing", return_tensors="pt")
11 outputs = model.generate(**inputs, max_length=100)
12 response = tokenizer.decode(outputs[0])
```

Listing 1: Basic Usage Example

## 6 Environmental Impact

### 6.1 Sustainability Metrics

- **Carbon Footprint:** 0.02 kg CO<sub>2</sub> per million inferences
- **Energy Usage:** 0.5 kWh per day (1000 users)
- **Efficiency Gain:** 98% reduction vs comparable models

### 6.2 Green AI Commitment

Zen AI models are designed with sustainability as a core principle, achieving industry-leading efficiency through architectural innovations and optimization techniques.

## 7 Safety and Alignment

### 7.1 Safety Measures

- Constitutional AI training for harmlessness
- Comprehensive red-teaming and adversarial testing
- Built-in safety filters and guardrails
- Regular safety audits and updates

### 7.2 Ethical Considerations

The model has been developed with careful attention to:

- Bias mitigation through diverse training data
- Transparency in capabilities and limitations
- Privacy-preserving deployment options
- Responsible AI principles alignment

## 8 Deployment Options

### 8.1 Available Formats

- **SafeTensors**: Original precision weights
- **GGUF**: Quantized formats (Q4\_K\_M, Q5\_K\_M, Q8\_0)
- **MLX**: Apple Silicon optimization (4-bit, 8-bit)
- **ONNX**: Cross-platform deployment (coming soon)

### 8.2 Hardware Requirements

| Precision | Memory | Recommended Hardware |
|-----------|--------|----------------------|
| FP16      | 1.2 GB | RTX 3060             |
| INT8      | 0.6 GB | GTX 1660             |
| INT4      | 2 GB   | Raspberry Pi 5       |

Table 4: Hardware Requirements by Precision

## 9 Future Work

### 9.1 Planned Improvements

- Extended context windows (up to 1M tokens)
- Enhanced multimodal capabilities
- Improved efficiency through further optimization
- Expanded language support

### 9.2 Research Directions

- Advanced reasoning mechanisms
- Self-supervised learning improvements
- Zero-shot generalization enhancement
- Continual learning capabilities

## 10 Conclusion

**Zen-Nano** represents a significant advancement in AI democratization, delivering exceptional performance for mobile/iot intelligence while maintaining unprecedented efficiency. Through innovative architecture design and careful optimization, the model achieves a balance between capability and sustainability that sets a new standard for responsible AI development.

## Acknowledgments

We thank the open-source community, our research partners, and the teams at Hanzo AI and Zoo Labs Foundation for their contributions to this work.

## References

### A Model Card

| Field         | Value   |
|---------------|---|
| Model Name    | Zen-Nano  |
| Version       | 1.0.0   |
| Release Date  | September 2025  |
| License       | Apache 2.0  |
| Repository    | <a href="https://huggingface.co/zenlm/zen-nano-0.6b-instruct">huggingface.co/zenlm/zen-nano-0.6b-instruct</a> |
| Documentation | <a href="https://github.com/zenlm/zen">github.com/zenlm/zen</a>   |
| Contact       | research@hanzo.ai   |

Table 5: Model Card Information