**Zen AI Model Family**

# Zen-Next

Flagship Model

Technical Whitepaper v1.0

Zach Kelling*
research@hanzo.ai

Zoo Labs Foundation
foundation@zoolabs.org

September 2025

**Abstract**

We present **Zen-Next**, a 80B parameter model optimized for flagship model. Built upon zen-72B, this model achieves state-of-the-art performance while maintaining exceptional efficiency with only 80B active parameters. Supporting 1M thinking tokens for advanced reasoning, the model represents a significant advancement in democratizing AI through sustainable and efficient architectures.

## Contents

**1 Introduction** **3**
1.1 Key Innovations . . . 3

**2 Architecture** **3**
2.1 Model Design . . . 3
2.2 Technical Innovations . . . 3
2.2.1 Mixture of Experts (MoE) . . . 3
2.2.2 Attention Mechanism . . . 3
2.2.3 Thinking Mode . . . 3

**3 Performance Benchmarks** **4**
3.1 Evaluation Results . . . 4
3.2 Efficiency Metrics . . . 4

**4 Training Methodology** **4**
4.1 Dataset . . . 4
4.2 Training Process . . . 4

**5 Use Cases and Applications** **4**
5.1 Primary Applications . . . 4
5.2 Integration Examples . . . 5

**6 Environmental Impact** **5**
6.1 Sustainability Metrics . . . 5
6.2 Green AI Commitment . . . 5

*zach@lux.network

**7 Safety and Alignment** **5**
7.1 Safety Measures . . . . . 5
7.2 Ethical Considerations . . . . . 5

**8 Deployment Options** **6**
8.1 Available Formats . . . . . 6
8.2 Hardware Requirements . . . . . 6

**9 Future Work** **6**
9.1 Planned Improvements . . . . . 6
9.2 Research Directions . . . . . 6

**10 Conclusion** **6**

**A Model Card** **7**

# 1 Introduction

The rapid advancement of artificial intelligence has created an unprecedented demand for models that balance capability with efficiency. **Zen-Next** addresses this challenge by delivering enterprise-grade performance while maintaining a minimal computational footprint.

## 1.1 Key Innovations

- **Efficient Architecture**: 80B active parameters from 80B total
- **Specialized Training**: Optimized for flagship model
- **Extended Context**: 128K context window
- **Thinking Mode**: 1M thinking tokens

# 2 Architecture

## 2.1 Model Design

Zen-Next is based on the zen-72B architecture with several key modifications:

| Component | Specification |
|---|---|
| Total Parameters | 80B |
| Active Parameters | 80B |
| Base Model | zen-72B |
| Context Length | 128K |
| Thinking Tokens | 1M |
| Architecture Type | Transformer |

Table 1: Zen-Next Architecture Specifications

## 2.2 Technical Innovations

### 2.2.1 Mixture of Experts (MoE)

The model uses a dense architecture with all parameters active during inference, optimized for maximum performance per parameter.

### 2.2.2 Attention Mechanism

Extended attention mechanisms support up to 128K context length with efficient KV-cache management.

### 2.2.3 Thinking Mode

Advanced reasoning through extended thinking tokens (up to 1M), enabling:

- Step-by-step problem decomposition
- Self-correction and verification
- Complex multi-step reasoning
- Internal deliberation before response

# 3 Performance Benchmarks

## 3.1 Evaluation Results

| Benchmark | Score |
|---|---|
| MMLU | 75.6% |
| HumanEval | 61.7% |
| GSM8K | 90.7% |
| HellaSwag | 86.9% |

Table 2: Language Understanding Benchmarks

## 3.2 Efficiency Metrics

| Metric | Value |
|---|---|
| Inference Speed | 45 tokens/sec |
| Memory Usage (INT4) | 40 GB |
| Energy Efficiency | 80% reduction |
| Latency (First Token) | 200 ms |

Table 3: Efficiency Metrics

# 4 Training Methodology

## 4.1 Dataset

The model was trained on a carefully curated dataset comprising:

- High-quality filtered web data (35TB)
- Domain-specific corpora for flagship model
- Synthetic data generation for edge cases
- Human feedback through RLHF

## 4.2 Training Process

1. **Pretraining**: 5 trillion tokens over 45 days on 64x A100
2. **Supervised Fine-tuning**: Task-specific optimization
3. **RLHF**: Alignment with human preferences
4. **Constitutional AI**: Safety and helpfulness optimization

# 5 Use Cases and Applications

## 5.1 Primary Applications

Conversational AI and chatbots

Content generation and summarization

Code completion and review

Educational assistance

Research and analysis

### 5.2 Integration Examples

```
from transformers import AutoModelForCausalLM, AutoTokenizer

# Load model and tokenizer
model = AutoModelForCausalLM.from_pretrained("zenlm/zen-next-80b-instruct")
tokenizer = AutoTokenizer.from_pretrained("zenlm/zen-next-80b-instruct")

# Generate response
inputs = tokenizer("Explain quantum computing", return_tensors="pt")
outputs = model.generate(**inputs, max_length=100)
response = tokenizer.decode(outputs[0])
```

Listing 1: Basic Usage Example

## 6 Environmental Impact

### 6.1 Sustainability Metrics

- **Carbon Footprint**: 0.45 kg COe per million inferences
- **Energy Usage**: 10.2 kWh per day (1000 users)
- **Efficiency Gain**: 80% reduction vs comparable models

### 6.2 Green AI Commitment

Zen AI models are designed with sustainability as a core principle, achieving industry-leading efficiency through architectural innovations and optimization techniques.

## 7 Safety and Alignment

### 7.1 Safety Measures

- Constitutional AI training for harmlessness
- Comprehensive red-teaming and adversarial testing
- Built-in safety filters and guardrails
- Regular safety audits and updates

### 7.2 Ethical Considerations

The model has been developed with careful attention to:

- Bias mitigation through diverse training data
- Transparency in capabilities and limitations
- Privacy-preserving deployment options
- Responsible AI principles alignment

# 8 Deployment Options

## 8.1 Available Formats

- **SafeTensors**: Original precision weights
- **GGUF**: Quantized formats (Q4_K_M, Q5_K_M, Q8_0)
- **MLX**: Apple Silicon optimization (4-bit, 8-bit)
- **ONNX**: Cross-platform deployment (coming soon)

## 8.2 Hardware Requirements

| Precision | Memory | Recommended Hardware |
|---|---|---|
| FP16 | 160 GB | 2x A100 80GB |
| INT8 | 80 GB | A100 80GB |
| INT4 | 40 GB | RTX 4090 |

Table 4: Hardware Requirements by Precision

# 9 Future Work

## 9.1 Planned Improvements

- Extended context windows (up to 1M tokens)
- Enhanced multimodal capabilities
- Improved efficiency through further optimization
- Expanded language support

## 9.2 Research Directions

- Advanced reasoning mechanisms
- Self-supervised learning improvements
- Zero-shot generalization enhancement
- Continual learning capabilities

# 10 Conclusion

**Zen-Next** represents a significant advancement in AI democratization, delivering exceptional performance for flagship model while maintaining unprecedented efficiency. Through innovative architecture design and careful optimization, the model achieves a balance between capability and sustainability that sets a new standard for responsible AI development.

# Acknowledgments

We thank the open-source community, our research partners, and the teams at Hanzo AI and Zoo Labs Foundation for their contributions to this work.

## References

# A Model Card

| Field | Value |
|---|---|
| Model Name | Zen-Next |
| Version | 1.0.0 |
| Release Date | September 2025 |
| License | Apache 2.0 |
| Repository | huggingface.co/zenlm/zen-next-80b-instruct |
| Documentation | github.com/zenlm/zen |
| Contact | research@hanzo.ai |

Table 5: Model Card Information