

Zen AI Model Family

Zen-3D

3D Scene Understanding and Generation at Scale

Technical Report v2025.01

Hanzo AI Research Team*

Zoo Labs Foundation[†]

January 2025

Abstract

We present **Zen-3D**, an 8-billion parameter model specialized for 3D scene understanding, mesh generation, and point cloud processing. Built on the Zen MoDE (Mixture of Distilled Experts) architecture with domain-adapted 3D tokenization, Zen-3D achieves state-of-the-art results across all major 3D benchmarks: 92.2% mean IoU on ShapeNet, 88.3% mIoU on ScanNet semantic segmentation, and 94.1% classification accuracy on ModelNet40. Unlike prior work that treats 3D understanding as a downstream task grafted onto language models, Zen-3D natively encodes spatial geometry through a novel triplane-based 3D tokenizer and dedicated attention mechanisms for volumetric context. The model supports interactive scene reconstruction, counterfactual 3D editing, and zero-shot generalization to unseen object categories—enabling downstream applications in robotics, augmented reality, and generative design at production scale.

Contents

1	Introduction	3
1.1	Model Overview	3
2	Architecture	3
2.1	Triplane 3D Tokenizer	3
2.2	Geometry-Aware Attention	4
2.3	Zen MoDE Backbone	4
2.4	Task-Specific Decoding Heads	4
3	Training	5
3.1	Dataset	5
3.2	Training Protocol	5
3.3	Data Augmentation	5
4	Evaluation	5
4.1	ShapeNet Object Reconstruction	5
4.2	ScanNet Semantic Segmentation	6
4.3	ModelNet40 Classification	6
4.4	Text-to-3D Generation Quality	6
4.5	Zero-Shot Generalization	7

*research@hanzo.ai

[†]foundation@zoo.ngo

5	Applications	7
5.1	Robotics Scene Understanding	7
5.2	Augmented Reality	7
5.3	Generative 3D Design	7
5.4	Scene Counterfactual Editing	7
6	Inference and Deployment	8
6.1	Inference Performance	8
6.2	Quantization	8
6.3	Python API	8
7	Related Work	8
7.1	3D Representation Learning	8
7.2	3D Generation	9
7.3	Vision-Language 3D Models	9
8	Conclusion	9

1 Introduction

Three-dimensional scene understanding has long been a frontier capability distinguishing physically situated AI from purely linguistic systems. Tasks such as reconstructing a scene from partial observations, segmenting semantically meaningful regions in a point cloud, or generating plausible 3D meshes from text descriptions require models that reason natively about geometry, topology, and spatial relations—not merely pattern-match against 2D projections.

Prior multimodal approaches to 3D have taken two broad strategies: (1) projecting 3D data into sequences of 2D renders and applying vision-language models, or (2) building bespoke 3D encoders that feed compressed representations into a language backbone. Both strategies suffer from either information loss in the projection step or a semantic gap between the 3D encoder and the pretrained language prior.

Zen-3D resolves this tension through three contributions:

1. **Native 3D tokenization:** A triplane decomposition tokenizer that converts point clouds and voxel grids into discrete token sequences without lossy projection.
2. **Geometry-aware attention:** Modified self-attention with relative 3D position encodings, enabling the model to reason about spatial adjacency, occlusion, and scale.
3. **Unified generation and understanding:** A single model that handles segmentation, reconstruction, and conditional mesh generation through task-specific decoding heads built on the shared Zen MoDE backbone.

1.1 Model Overview

Property	Value
Parameters	8B
Architecture	Zen MoDE (Mixture of Distilled Experts)
3D Tokenizer	Triplane VQ-VAE (codebook size 16,384)
Context Length	32,768 tokens (spatial + language)
Max Resolution	256^3 voxels / 100K point clouds
Output Modalities	Mesh, point cloud, segmentation mask, text
Training Data	47M 3D object-text pairs, 2.1M scene reconstructions

Table 1: Zen-3D Model Specifications

2 Architecture

2.1 Triplane 3D Tokenizer

Input 3D data—whether a point cloud, a voxel grid, or an implicit surface—is first projected onto three orthogonal feature planes (XY , XZ , YZ). Each plane is independently encoded by a convolutional backbone and then vector-quantized to produce a sequence of discrete tokens. The three plane token sequences are concatenated to form the full 3D token sequence fed to the transformer.

Formally, given input geometry $\mathcal{G} \in \mathbb{R}^{N \times 3}$:

$$F_{XY}, F_{XZ}, F_{YZ} = \text{TriplaneProject}(\mathcal{G}) \tag{1}$$

$$T_i = \text{VQ}(\text{ConvEnc}(F_i)), \quad i \in \{XY, XZ, YZ\} \tag{2}$$

$$T_{3D} = [T_{XY} \parallel T_{XZ} \parallel T_{YZ}] \tag{3}$$

The VQ-VAE codebook is trained jointly with the transformer backbone using a commitment loss $\mathcal{L}_{commit} = \|z_e - \text{sg}(z_q)\|_2^2$ where z_e is the encoder output and z_q is the nearest codebook entry.

2.2 Geometry-Aware Attention

Standard self-attention treats tokens as an unordered set, discarding geometric structure. We introduce **Spatial Relative Position Encoding (SRPE)**: for any pair of 3D tokens (i, j) , we compute their approximate 3D centroid distance d_{ij} and inject a learned distance bias into the attention logits:

$$a_{ij} = \frac{q_i \cdot k_j}{\sqrt{d_k}} + b(d_{ij}) \quad (4)$$

where $b : \mathbb{R}_{\geq 0} \rightarrow \mathbb{R}$ is a learned monotone function discretized into 64 distance buckets covering 0–10 meters in log-scale.

2.3 Zen MoDE Backbone

The transformer backbone follows the Zen MoDE architecture: 32 transformer layers, each with 16 attention heads and a mixture-of-experts feed-forward network with 8 experts (top-2 routing per token). The 3D token stream and optional text tokens are interleaved via cross-attention at every fourth layer, enabling bidirectional grounding between language and geometry.

Component	Configuration
Transformer layers	32
Attention heads	16
Hidden dimension	4096
FFN type	MoE (8 experts, top-2)
Active parameters per token	$\approx 2.1\text{B}$ of 8B total
Cross-attention layers	Every 4th (8 total)
Positional encoding	Rotary (RoPE) + SRPE

Table 2: Zen-3D Backbone Architecture

2.4 Task-Specific Decoding Heads

Three decoding heads share the backbone:

Segmentation head: A lightweight MLP applied to per-token backbone features, producing per-point semantic logits over a 200-class ontology (ShapeNet taxonomy + ScanNet classes).

Reconstruction head: A triplane decoder that inverts the tokenizer, predicting an occupancy field from which meshes are extracted via marching cubes at configurable resolution.

Generation head: A diffusion decoder conditioned on text-derived embeddings, producing novel 3D shapes by iteratively denoising a triplane latent.

3 Training

3.1 Dataset

Source	Samples	Proportion	Modality
ShapeNet v2	51,300	0.1%	Mesh + text
Objaverse 1.0	800,000	1.7%	Mesh + text
Objaverse-XL	10,200,000	21.6%	Mesh + text
ScanNet v2	1,513 scenes	0.003%	Point cloud + segmentation
3D-FRONT	18,000 scenes	0.04%	Scene mesh + layout
Synthetic augmentation	36,000,000	76.5%	Rendered point clouds
Total	47,070,813	100%	

Table 3: Zen-3D Training Data Composition

3.2 Training Protocol

Training proceeds in three stages:

Stage 1 – Tokenizer pretraining (50K steps): The VQ-VAE is trained standalone on the full 3D corpus with reconstruction loss and commitment loss. Codebook utilization is monitored; dead entries are re-initialized via exponential moving average updates.

Stage 2 – Backbone pretraining (200K steps): The Zen MoDE backbone is pretrained on tokenized 3D sequences with a masked-token prediction objective, analogous to masked language modeling but over the 3D token space.

Stage 3 – Task fine-tuning (100K steps): All three decoding heads are trained jointly with task-specific losses: cross-entropy for segmentation, binary cross-entropy occupancy loss for reconstruction, and diffusion DDPM loss for generation. Tasks are sampled at ratio 4:3:3.

Stage	Steps	Batch	LR	Hardware
Tokenizer	50K	512	1e-4	16×A100 80GB
Backbone	200K	256	3e-4	64×A100 80GB
Fine-tune	100K	128	5e-5	32×A100 80GB

Table 4: Training Configuration by Stage

3.3 Data Augmentation

To improve robustness to real-world sensor noise, all point cloud inputs during training undergo: random dropout (10–30% of points), Gaussian jitter ($\sigma = 0.01$ m), random rotation about the vertical axis, and random scale perturbation ($\pm 10\%$).

4 Evaluation

4.1 ShapeNet Object Reconstruction

We evaluate 3D reconstruction quality on the ShapeNet v2 test split using Intersection over Union (IoU) at 0.5 occupancy threshold.

Model	Chair	Table	Lamp	Car	Mean IoU
ONet	82.3	79.1	68.4	83.7	75.9
ConvONet	88.6	85.2	74.1	88.9	82.4
IF-Net	90.1	87.3	76.8	90.2	84.6
3D-LLM	91.4	88.7	78.3	91.5	87.2
Zen-3D	95.2	93.8	84.1	95.7	92.2

Table 5: ShapeNet Reconstruction (IoU %, higher is better)

4.2 ScanNet Semantic Segmentation

Point cloud semantic segmentation on ScanNet v2 validation set (20 classes).

Model	mIoU	mAcc	oAcc	Parameters
PointNet++	53.5	64.2	85.3	1.5M
KPConv	68.4	79.1	88.1	14.3M
PointTransformer	70.6	81.9	90.2	7.8M
Mask3D	73.7	83.4	91.0	38.4M
Point-MAE + ScanNet	80.3	87.2	91.8	22.1M
Zen-3D	88.3	92.1	94.6	8B

Table 6: ScanNet Semantic Segmentation (validation, higher is better)

4.3 ModelNet40 Classification

Object classification on ModelNet40 with 1,024 input points, 10-fold cross-validation.

Model	OA (%)	mAcc (%)
PointNet	89.2	86.2
PointNet++	91.9	90.7
DGCNN	92.9	90.2
Point-BERT	93.2	90.6
Point-MAE	93.8	90.6
PointGPT-L	94.0	91.1
Zen-3D	94.1	91.3

Table 7: ModelNet40 Classification (OA = Overall Accuracy)

4.4 Text-to-3D Generation Quality

We evaluate generative quality using the Frechet Point Cloud Distance (FPCD) and CLIP-3D alignment score on a held-out set of 5,000 text prompts from the Cap3D dataset.

Model	FPCD ↓	CLIP-3D ↑
Shap-E	143.2	0.312
Point-E	128.7	0.331
DreamFusion	115.4	0.347
One-2-3-45	98.3	0.369
Zen-3D	71.4	0.403

Table 8: Text-to-3D Generation (Cap3D evaluation set)

4.5 Zero-Shot Generalization

To assess zero-shot generalization, we evaluate on the Objaverse-LVIS benchmark (1,156 novel categories not seen during training). Zen-3D achieves 61.3% top-1 classification accuracy, compared to 43.7% for the best prior approach, demonstrating strong zero-shot transfer from the large-scale pretraining regime.

5 Applications

5.1 Robotics Scene Understanding

Zen-3D has been integrated into the Hanzo robotics perception stack, where it operates on streaming LiDAR point clouds at 5 Hz (200ms per frame on a single A10G GPU). The model simultaneously provides semantic segmentation for grasp planning and 3D occupancy maps for collision-free path planning.

5.2 Augmented Reality

For AR applications, Zen-3D’s reconstruction head is used to generate watertight meshes from partial RGBD observations in real time, enabling stable AR object placement without depth sensor drift artifacts.

5.3 Generative 3D Design

The text-to-3D generation pipeline enables designers to specify 3D shapes in natural language and receive printable mesh outputs in under 10 seconds. Integration with standard CAD tools (FreeCAD, Blender) is provided via the Zen SDK.

5.4 Scene Counterfactual Editing

Zen-3D supports counterfactual scene editing: given a complete scene reconstruction and a natural language instruction (e.g., “remove the chair and place a desk”), the model regenerates the affected scene region while maintaining geometric consistency with the unchanged context.

6 Inference and Deployment

6.1 Inference Performance

Task	Latency (ms)	Throughput	Hardware
Segmentation (10K pts)	45	22 scenes/s	A10G
Reconstruction (ShapeNet)	120	8 meshes/s	A10G
Text-to-3D generation	8,200	0.12 shapes/s	A100
Classification (1K pts)	12	83 objects/s	T4

Table 9: Zen-3D Inference Performance

6.2 Quantization

Zen-3D supports INT8 quantization via the Zen SDK with less than 0.5% degradation on all benchmarks. The quantized model fits in 10GB VRAM, enabling deployment on consumer-grade GPUs (RTX 3080 and above).

6.3 Python API

```
1 from zen import Zen3D
2 import numpy as np
3
4 model = Zen3D.from_pretrained("zenlm/zen-3d-8b")
5
6 # Semantic segmentation from point cloud
7 points = np.load("scene.npy") # (N, 3) float32
8 labels = model.segment(points, taxonomy="scannet20")
9
10 # Text-to-3D generation
11 mesh = model.generate(
12     "a modern office chair with armrests",
13     resolution=128,
14     num_steps=50
15 )
16 mesh.export("chair.obj")
17
18 # Scene reconstruction from partial observations
19 partial = points[:points.shape[0] // 2]
20 complete = model.reconstruct(partial)
```

Listing 1: Zen-3D Inference Example

7 Related Work

7.1 3D Representation Learning

PointNet [1] established the paradigm of operating directly on unordered point sets via symmetric functions. PointNet++ [2] extended this with hierarchical local feature learning. Transformer-based approaches (Point-BERT, Point-MAE) adapted masked pretraining from NLP to 3D, achieving strong transfer performance.

7.2 3D Generation

Occupancy Networks [3] and Implicit Functions [4] enabled continuous implicit representations. Diffusion-based generators (Shap-E, Point-E) brought generative quality closer to 2D image synthesis. Zen-3D advances this line by conditioning generation on the same backbone used for understanding, enabling tighter semantic alignment.

7.3 Vision-Language 3D Models

3D-LLM and related work demonstrated that language model priors can improve 3D understanding. Zen-3D builds on this direction but replaces the 2D-projection interface with native 3D tokenization, eliminating the projection bottleneck.

8 Conclusion

Zen-3D demonstrates that a single 8B parameter model, trained with the Zen MoDE architecture and a purpose-built triplane tokenizer, can simultaneously achieve state-of-the-art results across 3D segmentation (ScanNet mIoU 88.3%), reconstruction (ShapeNet mean IoU 92.2%), and classification (ModelNet40 94.1%). The unified backbone enables cross-task transfer that benefits all three capabilities simultaneously.

Future work will extend Zen-3D to dynamic 4D scene understanding (point cloud video), tighter integration with the Zen-Omni multimodal backbone for joint 2D+3D understanding, and distillation of the 8B model to a sub-1B edge variant for real-time mobile AR applications.

References

- [1] R. Q. Charles et al., “PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation,” CVPR, 2017.
- [2] C. R. Qi et al., “PointNet++: Deep Hierarchical Feature Learning on Point Sets in a Metric Space,” NeurIPS, 2017.
- [3] L. Mescheder et al., “Occupancy Networks: Learning 3D Reconstruction in Function Space,” CVPR, 2019.
- [4] J. J. Park et al., “DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation,” CVPR, 2019.