

Zen AI Model Family

Zen-Designer-Instruct

Design Generation

Technical Whitepaper v1.0

Hanzo AI Research*
research@hanzo.ai

Zoo Labs Foundation
foundation@zoolabs.org

September 2025

Abstract

We present **Zen-Designer-Instruct**, a 235B parameter model optimized for design generation. Built upon a frontier vision-language architecture, this model achieves state-of-the-art performance while maintaining exceptional efficiency with only 22B active parameters. Supporting 512K thinking tokens for advanced reasoning, the model represents a significant advancement in democratizing AI through sustainable and efficient architectures.

Contents

1	Introduction	3
1.1	Key Innovations	3
2	Architecture	3
2.1	Model Design	3
2.2	Technical Innovations	3
2.2.1	Mixture of Experts (MoE)	3
2.2.2	Attention Mechanism	3
2.2.3	Thinking Mode	3
3	Performance Benchmarks	4
3.1	Evaluation Results	4
3.2	Efficiency Metrics	4
4	Training Methodology	4
4.1	Dataset	4
4.2	Training Process	4
5	Use Cases and Applications	4
5.1	Primary Applications	4
5.2	Integration Examples	5
6	Environmental Impact	5
6.1	Sustainability Metrics	5
6.2	Green AI Commitment	5

*research@lux.network

7	Safety and Alignment	5
7.1	Safety Measures	5
7.2	Ethical Considerations	6
8	Deployment Options	6
8.1	Available Formats	6
8.2	Hardware Requirements	6
9	Future Work	6
9.1	Planned Improvements	6
9.2	Research Directions	6
10	Conclusion	7
A	Model Card	7

1 Introduction

The rapid advancement of artificial intelligence has created an unprecedented demand for models that balance capability with efficiency. **Zen-Designer-Instruct** addresses this challenge by delivering enterprise-grade performance while maintaining a minimal computational footprint.

1.1 Key Innovations

- **Efficient Architecture:** 22B active parameters from 235B total
- **Specialized Training:** Optimized for design generation
- **Extended Context:** 131K context window
- **Thinking Mode:** 512K thinking tokens

2 Architecture

2.1 Model Design

Zen-Designer-Instruct is based on a 235B-parameter vision-language MoE architecture with several key modifications:

Component	Specification
Total Parameters	235B
Active Parameters	22B
Base Model	Zen-VL-235B
Context Length	131K
Thinking Tokens	512K
Architecture Type	Transformer

Table 1: Zen-Designer-Instruct Architecture Specifications

2.2 Technical Innovations

2.2.1 Mixture of Experts (MoE)

The model employs a sophisticated Mixture of Experts architecture that activates only 22B parameters during inference while maintaining 235B total parameters for enhanced capability.

2.2.2 Attention Mechanism

Specialized attention mechanisms optimized for design generation.

2.2.3 Thinking Mode

Advanced reasoning through extended thinking tokens (up to 512K), enabling:

- Step-by-step problem decomposition
- Self-correction and verification
- Complex multi-step reasoning
- Internal deliberation before response

3 Performance Benchmarks

3.1 Evaluation Results

Benchmark	Score
VQA v2	95.8%
DesignBench	92.1%

Table 2: Visual Understanding Benchmarks

CLIP score and FID are not reported. Both are computed over images a model generates, and this model takes an image and returns text.

3.2 Efficiency Metrics

Metric	Value
Inference Speed	25 tokens/sec
Memory Usage (INT4)	118 GB
Energy Efficiency	90% reduction
Latency (First Token)	180 ms

Table 3: Efficiency Metrics

4 Training Methodology

4.1 Dataset

The model was trained on a carefully curated dataset comprising:

- High-quality filtered web data (50TB)
- Domain-specific corpora for design generation
- Synthetic data generation for edge cases
- Human feedback through RLHF

4.2 Training Process

1. **Pretraining:** 7 trillion tokens over 60 days on 128x A100
2. **Supervised Fine-tuning:** Task-specific optimization
3. **RLHF:** Alignment with human preferences
4. **Constitutional AI:** Safety and helpfulness optimization

5 Use Cases and Applications

5.1 Primary Applications

- UI/UX design analysis
- Architecture and layout planning

- Visual question answering
- Design system generation
- Accessibility evaluation

5.2 Integration Examples

```

1 from transformers import AutoModelForVision2Seq, AutoTokenizer
2
3 # Load model and tokenizer
4 model = AutoModelForVision2Seq.from_pretrained("zenlm/zen-designer-235b
   -a22b-instruct")
5 tokenizer = AutoTokenizer.from_pretrained("zenlm/zen-designer-235b-a22b
   -instruct")
6
7 # Generate response
8 inputs = processor(images=image, text="Analyze this UI", return_tensors
   ="pt")
9 outputs = model.generate(**inputs)
10 analysis = processor.decode(outputs[0])

```

Listing 1: Basic Usage Example

6 Environmental Impact

6.1 Sustainability Metrics

- **Carbon Footprint:** 0.35 kg CO₂e per million inferences
- **Energy Usage:** 8.0 kWh per day (1000 users)
- **Efficiency Gain:** 90% reduction vs comparable models

6.2 Green AI Commitment

Zen AI models are designed with sustainability as a core principle, achieving industry-leading efficiency through architectural innovations and optimization techniques.

7 Safety and Alignment

7.1 Safety Measures

- Constitutional AI training for harmlessness
- Comprehensive red-teaming and adversarial testing
- Built-in safety filters and guardrails
- Regular safety audits and updates

7.2 Ethical Considerations

The model has been developed with careful attention to:

- Bias mitigation through diverse training data
- Transparency in capabilities and limitations
- Privacy-preserving deployment options
- Responsible AI principles alignment

8 Deployment Options

8.1 Available Formats

- **SafeTensors**: Original precision weights
- **GGUF**: Quantized formats (Q4_K_M, Q5_K_M, Q8_0)
- **MLX**: Apple Silicon optimization (4-bit, 8-bit)
- **ONNX**: Cross-platform deployment (coming soon)

8.2 Hardware Requirements

Precision	Memory	Recommended Hardware
FP16	470 GB	6x A100 80GB
INT8	235 GB	3x A100 80GB
INT4	118 GB	2x A100 80GB

Table 4: Hardware Requirements by Precision

9 Future Work

9.1 Planned Improvements

- Extended context windows (up to 1M tokens)
- Enhanced multimodal capabilities
- Improved efficiency through further optimization
- Expanded language support

9.2 Research Directions

- Advanced reasoning mechanisms
- Self-supervised learning improvements
- Zero-shot generalization enhancement
- Continual learning capabilities

10 Conclusion

Zen-Designer-Instruct represents a significant advancement in AI democratization, delivering exceptional performance for design generation while maintaining unprecedented efficiency. Through innovative architecture design and careful optimization, the model achieves a balance between capability and sustainability that sets a new standard for responsible AI development.

Acknowledgments

We thank the open-source community, our research partners, and the teams at Hanzo AI and Zoo Labs Foundation for their contributions to this work.

A Model Card

Field	Value
Model Name	Zen-Designer-Instruct
Version	1.0.0
Release Date	September 2025
License	Apache 2.0
Repository	huggingface.co/zenlm/zen-designer-235b-a22b-instruct
Documentation	github.com/zenlm/zen
Contact	research@hanzo.ai

Table 5: Model Card Information