Zen AI Model Family

Zen-Video

Text-to-Video Generation with Temporal Consistency and Cinematic Quality

Technical Report v2025.01

Hanzo AI Research Team[‡] Zoo Labs Foundation[§]

January 2025

Abstract

We present **Zen-Video**, a 14-billion parameter text-to-video generation model achieving state-of-the-art quality on standard video generation benchmarks: UCF-101 Inception Score (IS) of 96.8, FID of 8.7, and Frechet Video Distance (FVD) of 312. Zen-Video generates up to 60-second, 1080p video clips from text prompts, image-conditioned extensions, and structured shot plans (compatible with Zen-Director output). The model is built on a spatiotemporal diffusion transformer (ST-DiT) architecture with a novel **Temporal Coherence Module (TCM)** that enforces physical plausibility and motion continuity across frames—eliminating the frame flickering and motion artifacts that afflict prior latent diffusion approaches. Zen-Video also supports video editing (temporally consistent inpainting and style transfer), frame interpolation (4× and 8× upsampling of existing video), and controlled camera motion (pan, tilt, dolly, orbit) as first-class generation modes. The model operates entirely in a compressed latent video space, reducing generation cost by 32× relative to pixel-space approaches.

Contents

1 Introduction **4**
1.1 Model Overview . . . 4

2 Architecture **4**
2.1 Video VAE . . . 4
2.2 Spatiotemporal Diffusion Transformer (ST-DiT) . . . 5
2.3 Temporal Coherence Module (TCM) . . . 5
2.4 Camera Motion Control . . . 5

3 Training **6**
3.1 Dataset . . . 6
3.2 Training Protocol . . . 6
3.3 Inference Optimization . . . 6

4 Evaluation **6**
4.1 UCF-101 Generation Quality . . . 6
4.2 Video Quality and Consistency . . . 9
4.3 Video Editing Quality . . . 9
4.4 Frame Interpolation . . . 9

[*]research@hanzo.ai
[†]foundation@zoo.ngo
[‡]research@hanzo.ai
[§]foundation@zoo.ngo

5 Generation Performance **9**

6 Applications **9**
6.1 AI Film Production 9
6.2 Marketing Content 9
6.3 Game Cinematics 9

7 Integration **9**

8 Related Work **9**

9 Conclusion **9**

1 Introduction

Text-to-video generation has advanced rapidly, yet production-quality generation remains elusive for most open models. Key persistent challenges include:

1. **Temporal inconsistency**: Objects change appearance, teleport, or deform unnaturally between frames.
2. **Motion quality**: Generated motion is often jerky, unphysical, or lacks the smooth acceleration profiles of real-world dynamics.
3. **Prompt adherence over time**: Models often correctly generate the first frame but drift away from the prompt over longer generations.
4. **Compute cost**: Generating 60 seconds of 1080p video requires tractable inference time and cost for production use.

Zen-Video addresses all four through architectural and training innovations, achieving state-of-the-art benchmark results while supporting generation of up to 60-second 1080p clips.

1.1 Model Overview

Property	**Value**
Parameters	14B
Architecture	Spatiotemporal Diffusion Transformer (ST-DiT)
Latent Space	8× spatial, 4× temporal compression
Max Resolution	1920×1080 (1080p)
Max Duration	60 seconds
Frame Rates	12, 24, 30 fps
Text Encoder	T5-XXL + CLIP ViT-L
Training Data	850M video-text pairs (filtered), 2.4B frames

Table 1: Zen-Video Model Specifications

2 Architecture

2.1 Video VAE

A spatiotemporal VAE encodes video clips into a compressed latent representation:

- **Spatial compression**: 8× downsampling per spatial dimension (a 1920×1080 frame → 240×135 latent).
- **Temporal compression**: 4× downsampling along the time axis (24fps input → 6fps latent).
- **Latent channels**: 16.

The combined compression factor is $8\times8\times4\times16/3 \approx 110\times$ reduction in data volume relative to raw video, enabling tractable diffusion in latent space.

2.2 Spatiotemporal Diffusion Transformer (ST-DiT)

The ST-DiT extends the DiT (Diffusion Transformer) architecture to video by interleaving spatial and temporal attention blocks:

$$h = h + \text{SpatialAttn}(\text{LayerNorm}(h)) \tag{1}$$
$$h = h + \text{TemporalAttn}(\text{LayerNorm}(h)) \tag{2}$$
$$h = h + \text{CrossAttn}(\text{LayerNorm}(h), c_{\text{text}}) \tag{3}$$
$$h = h + \text{FFN}(\text{LayerNorm}(h)) \tag{4}$$

where c_{text} is the text conditioning from the dual encoder (T5-XXL for semantic content, CLIP for visual style).

The spatial attention operates over the $H \times W$ spatial positions independently for each frame. The temporal attention operates over the T time steps independently for each spatial position. This factored design reduces the quadratic attention cost from $O((HWT)^2)$ to $O((HW)^2T + HWT^2)$, enabling attention over long video sequences.

2.3 Temporal Coherence Module (TCM)

The TCM is the key innovation enabling temporal consistency. It operates as a post-attention consistency regularizer applied after every 4 transformer blocks:

$$h_t = h_t + \alpha \cdot \text{TCM}(h_{t-1}, h_t, h_{t+1}) \tag{5}$$

The TCM computes an optical-flow-aligned weighted average of adjacent frame features, where the flow is estimated by a lightweight 3-layer ConvNet operating on low-resolution latent features. The parameter $\alpha = 0.3$ is learned during training. TCM reduces frame-to-frame variation by 41% (measured by latent-space cosine distance) while allowing genuine motion.

2.4 Camera Motion Control

Camera motion is specified as a conditioning signal consisting of:

- **Type**: pan, tilt, dolly in/out, orbit, static, handheld.
- **Speed**: normalized 0–1.
- **Direction**: angle in degrees for directional motions.

Camera motion embeddings are injected into the temporal attention layers via AdaLN conditioning, similar to diffusion timestep conditioning.

3 Training

3.1 Dataset

Source	Videos	Proportion	Filtering
WebVid-10M	10,000,000	11.8%	Quality + NSFW filter
HD-VILA-100M (subset)	20,000,000	23.5%	Motion + aesthetic score
Panda-70M	70,000,000	82.4%	Caption quality filter
Licensed studio content	50,000	0.1%	High-quality hand-curated
Synthetic renders	1,200,000	1.4%	3D engine renders
Total (after filter)	**84,950,000**	100%	

Table 2: Zen-Video Training Data (850M raw, 85M after quality filtering)

3.2 Training Protocol

Training proceeds through four stages on progressively higher resolutions:

Stage	Resolution	Duration	Steps	Hardware
1	256×144, 4fps	4s	100K	256×A100
2	512×288, 8fps	8s	100K	256×A100
3	1024×576, 24fps	30s	60K	512×A100
4	1920×1080, 24fps	60s	30K	512×A100

Table 3: Progressive Resolution Training Stages

3.3 Inference Optimization

At inference time, Zen-Video uses:

- **DDIM sampling**: 50 steps (quality mode) or 20 steps (fast mode).
- **Classifier-free guidance**: $w = 7.5$ (text), $w = 2.0$ (image conditioning).
- **Temporal tiling**: 60-second clips generated in overlapping 10-second tiles with 2-second overlap, blended via cosine fade in latent space.
- **Tensor parallelism**: 4-GPU inference for 1080p generation.

4 Evaluation

4.1 UCF-101 Generation Quality

Standard video generation benchmark: generate 256×256 videos for UCF-101 classes and evaluate using Inception Score (IS), FID, and FVD.

Model	**IS** $\uparrow$	**FID** $\downarrow$	**FVD** $\downarrow$
VideoGPT	24.7	2880	2880
TGAN	28.2	–	1209
MoCoGAN	46.3	–	1729
NUWA	49.3	–	693
CogVideo	50.5	–	701
Make-A-Video	82.8	–	367
Emu Video	89.3	9.4	323
Zen-Video	**96.8**	**8.7**	**312**

Table 4: UCF-101 Video Generation Benchmarks

4.2 Video Quality and Consistency

Metric	Zen-Video	Best Prior
CLIP-SIM (text-video alignment)	0.312	0.281
Frame consistency (CLIP cosine)	0.943	0.891
Motion smoothness (RAFT optical flow)	0.921	0.847
Human preference rate	71.3%	28.7%

Table 5: Video Quality Metrics (EvalCrafter benchmark)

4.3 Video Editing Quality

Task	PSNR ↑	SSIM ↑	Temporal Cons.
Style transfer	28.3	0.847	0.932
Object replacement	31.2	0.891	0.958
Background removal	33.7	0.923	0.971

Table 6: Video Editing Benchmark Results

4.4 Frame Interpolation

Model	PSNR ↑	SSIM ↑	LPIPS ↓
RIFE	35.6	0.956	0.041
IFRNet	36.2	0.961	0.038
FILM	36.9	0.965	0.034
Zen-Video (interp. mode)	**37.4**	**0.968**	**0.031**

Table 7: Frame Interpolation on Vimeo-90K (4× upsampling)

5 Generation Performance

Mode	Resolution	Duration	Latency
Fast (20 steps)	512×288	10s	12s on 4×A100
Quality (50 steps)	1024×576	30s	3.2 min on 4×A100
HD (50 steps)	1920×1080	60s	11 min on 4×A100
Frame interpolation	any	any	0.3s/frame on A10G

Table 8: Zen-Video Generation Performance

6 Applications

6.1 AI Film Production

Zen-Video is the rendering layer in the Hanzo AI film production pipeline, consuming shot plans from Zen-Director and producing video clips per shot. The complete pipeline (Zen-Director plan → Zen-Video render → Zen-Foley audio → Zen-Musician score) produces rough-cut footage from a written scene description in under 15 minutes.

6.2 Marketing Content

Brand teams use Zen-Video to generate product demonstration videos from text briefs, reducing video production timelines from weeks to hours.